

Plexus Discovery

A specification for an autonomous mechanism-search campaign
on Okuda tubulation

Janelia Research Campus

July 31, 2026

Abstract

This is a build specification, not an essay. It defines an agentic loop that searches for the *mechanism* behind a morphogenetic phenotype — as a composition of Plexus operators — and that is designed to run unattended on a GPU partition for weeks without descending into the failure mode that consumed our last forty-four rounds.

The target system is Okuda tubulation: a closed epithelial shell that grows a tube. We already make a tube; we make it *by the wrong mechanism*, and we know this because relaxing our tube toward mechanical equilibrium destroys it (§1). That gives the campaign a sharp, falsifiable objective rather than a resemblance target.

Section 2 states what machinery already exists and transfers. Section 3 defines the Okuda composition space concretely — operators, ports, stages, legal edits. Section 4 diagnoses why searches like ours go into rabbit holes, from our own record with numbers. Section 5 is the core of the document: an explicit control law with a batch-triage-truncate-starve-freeze-terminate cycle, so that budget is *taken away* from unproductive branches automatically. Section 7 assigns the work to agents, borrowing orchestration from two multi-agent systems published together in *Nature* in July 2026 while deliberately diverging from their adjudication. Sections 11–12 define the artefacts, including the causal descriptions and the trigger that hands a problem to the differentiable fit. Section 13 is a pre-flight checklist of eleven defects that must be closed first — five of them found while writing this document.

Contents

1	The question the campaign answers	3
1.1	Not “can we make a tube”	3
1.2	The question is which mechanism holds it open	3
1.3	The campaign objective: a causal lever-map, not a target	3
1.4	The cells of the map, and coverage as progress	4
1.5	The metric that makes the objective measurable	4
2	What already exists and transfers	5
2.1	Transfers unchanged — the three structural guarantees	5
2.2	Must be replaced	6
2.3	Must be added	6
3	The Okuda composition space	6
3.1	Vocabulary	6
3.2	Seed and legal edits	7

3.3	Phenotype and the metric bank	7
4	Why searches like ours go into rabbit holes	7
5	The control law	8
5.1	One round	8
5.2	Termination and escalation	9
5.3	Rules enforced in code, not recommended in prose	9
6	The scientific protocol: hypothesis first	10
6.1	Two intents, and why the mixture matters	10
6.2	The 70/30 balance is a setpoint, not a quota	10
6.3	The knowledge document	11
7	The agents	11
8	Knowledge certification, and who may touch the substrate	13
8.1	A worked example of why the role must exist	13
8.2	The Metrologist	14
8.3	The boundary: who may modify the substrate	14
9	Which agents are language models, and which must not be	14
10	Running for weeks: where this breaks	16
10.1	Knowledge growth, as a measured quantity	16
10.2	The video battery	16
11	What the campaign produces	17
11.1	The ledger	17
11.2	Causal descriptions — “this spec gives this phenomenon”	17
11.3	The operator backlog — atlas growth	17
11.4	Movies	17
11.5	One differentiable figure	18
12	Where the differentiable version enters	18
13	Pre-flight: eleven defects that must be closed first	18
14	Operating it	20
15	Status and the switch-on gate	21
16	First live results (2026-07-30, overnight)	21
16.1	The composition mode could never have completed a round	21
16.2	The surprise rate was computed by a regex that could only bias it downwards	22
16.3	The supervisor did not survive a restart	22
16.4	The escalation path, and its first entry	22
16.5	Turing $\times$ vertex: a structural limitation, measured	22
16.6	What the defences caught, and what they cost me	23

17 External review — 2026-07-31, 07:58 EDT	24
18 Plan of record — 2026-07-31, 08:05 EDT	25
18.1 Three diagnoses that set the order of work	25
18.2 Okuda Table 1 versus the search box	25
18.3 Phases	26

1 The question the campaign answers

1.1 Not “can we make a tube”

We can. A tip-tracking activation driver on a growing, dividing vesicle produces a long, thin, constant-diameter tube with a length-to-diameter ratio above four, and its diameter tracks the morphogen diffusion coefficient — a qualitative reproduction of the published law. That question is closed.

1.2 The question is which mechanism holds it open

Round 41 of that campaign asked what happens if the tissue is pushed toward the published *regime* — Okuda’s model is explicitly quasi-static, relaxing to force balance between growth events. The answer was decisive:

Relaxation destroys our tube. More relaxation collapses it back toward a sphere. Our tube is held open by an explicit outward force; remove the force and the shape energy pulls it in. The published tube is an *equilibrium* shape produced by localised proliferation.

The mechanical signature agrees: force and pressure are concentrated entirely in the protrusion, the body is idle, and division and migration pile up at the same tip. That is what a forced protrusion looks like. A grown one shares load with the body.

So two mechanisms produce the same picture, and the picture is exactly what cannot tell them apart. This is the general statement of why parameter search cannot find a mechanism, and it is stronger than the usual one: *the search optimises resemblance, and resemblance is the quantity blind to the distinction that matters.*

1.3 The campaign objective: a causal lever-map, not a target

An earlier draft of this note stated the objective as one question — *produce a sustained tube as a growth-driven quasi-static equilibrium*. That framing is wrong, and the throughput arithmetic of §10 is what exposed it: with a single target the search *converges*, freezes every cluster within days, and stops, having established one fact. It also wastes most of its own evidence, because a composition that does not make a tube teaches nothing about tubes.

The objective is the one plexus2’s Loop I already names:

Build the CAUSAL LEVER-MAP of the mechanism space: for every operator, every implementation and every routing, what does it do alone, and what does it do in combination — since mixtures rarely surrender their causal structure to inspection.

The product is a map, not a winner. Specific questions become **queries against it**:

query	answered once the relevant cells are covered
sustained tube phase diagram	a tube that survives relaxation, <i>without</i> the extrusion forcing node Okuda’s four (χ, γ) morphologies — thin tube, thick tube, undulation, branching — from <i>one</i> composition
necessity impossibility	for each observed phenotype, which operators are load-bearing which composition families provably cannot reach a phenotype

1.4 The cells of the map, and coverage as progress

Solo	for each operator and each structural implementation, its effect alone
Pair	for each pair, is the joint effect what the two solo effects predict (<i>additive</i>), or not? Interactions are the expensive knowledge — they are precisely what cannot be read off the code
Routing	morphogen $\rightarrow$ growth.gate vs $\rightarrow$ divide.axis vs $\rightarrow$ extrude.site: same op- erators, different mechanism
Necessity	for each observed phenotype, which ablations destroy it
Phenotype	which regions produce which morphology — labels <i>discovered</i> by the analysts, never declared in advance

Progress is then **coverage**: the fraction of cells carrying enough evidence to state a verdict. Every run fills a cell, so no run is wasted, and the campaign has a natural extent — it runs until the levers are characterised. That is a weeks-scale programme rather than a days-scale search, and it is what makes the multi-week campaign coherent rather than merely long.

Two subordinate questions inherit from the record and enter the map as ordinary open cells rather than as blockers: whether patterning can be *emergent* (a late round concluded the activator–inhibitor system floods the growing domain and cannot maintain a tip-riding peak — an impossibility claim made *before* the precondition checks existed, so it re-enters as **Open**), and whether per-cell three-dimensional volume changes the answer (the monolayer implementation, whose first wiring spiked because division still triggered on the wedge volume rather than $A \cdot h$).

1.5 The metric that makes the objective measurable

The objective above is useless without an observable that separates a forced tube from a grown one. No metric in the existing bank does. We therefore define one, and it is cheap:

The quasi-static test. Take a run’s final state; continue the simulation with *growth and any driver switched off* and only mechanics and reconnection active; measure what fraction of the protrusion survives.

$$Q = \frac{\text{aspect after relaxation}}{\text{aspect before relaxation}}.$$

A forced tube collapses ($Q \rightarrow 0$). An equilibrium tube persists ($Q \rightarrow 1$). Alongside it we record the pressure ratio between tube and body, which should approach unity for a grown tube and is presently around three.

Q is the campaign’s primary discriminator. It is defined before the search starts, it is falsifiable, and it is the direct operationalisation of the round-41 finding. Authoring it is an instance of the discipline this whole document is built around: *when the eye sees a distinction no number captures, the missing number is the deliverable*.

2 What already exists and transfers

A working three-loop discovery system exists — roughly 1,100 lines — and has run unattended for six rounds on embryonic gland branching, producing 1,216 archived runs and one conclusion that a later, better measurement overturned. It is the substrate for this campaign. This section records what transfers unchanged, what must be replaced, and what must be added.

2.1 Transfers unchanged — the three structural guarantees

1. A composition is a typed graph, and its identity excludes parameters. A candidate mechanism is operator nodes plus typed connections between output ports and input slots. A morphogen may be routed to a cleft **source**, a growth **gate**, a chemotaxis **field** or an anisotropy **axis** — four *different mechanisms*, not four settings. Identity is a hash of the structure only:

```
def comp_hash(composition):  
    # operators + typed connections ONLY -- no parameter values, no seed
```

Therefore **a change of numbers provably cannot register as a new hypothesis**. This single line is the structural fix for what happened over rounds 1–30, and it is why the campaign can be trusted to count its own progress honestly.

2. Evidence and interpretation are separate objects. The rule is enforced in code, not in prose:

$$\text{simulation} \longrightarrow \text{RunRecord} \longrightarrow \text{knowledge} \quad (\text{never } \text{simulation} \longrightarrow \text{knowledge}).$$

A **RunRecord** carries the raw facts of one simulation — composition, parameters, seed, backend, initial condition, parent, and the single edit that produced it — and is *immutable*: assigning to a raw fact raises. Measurements attach as *versioned*, *append-only* analyses, so **metric_v2** sits beside **metric_v1** and neither is lost. The archive is append-only, and run identity is a content hash of exactly the reproducible inputs, so a crashed campaign resumes without duplicating or losing evidence.

This is what let a new observable, in the gland campaign, **re-score archived runs and change a mechanistic conclusion without re-simulating anything**. It is the most valuable property the system has.

3. One legal edit at a time, stage-gated. Four moves exist — add, remove, connect, disconnect — each returning a new graph together with the edit that produced it. Every run records its parent and its mutation, so the archive is a causal tree, not a bag of configurations. Stage gating keeps the legal move set small and interpretable rather than an enumeration of 2^N subsets.

2.2 Must be replaced

Component	Why it must change
Operator vocabulary	The existing one is phase-field gland vocabulary. Okuda needs its own (§3).
Frontier expansion	Breadth-first with a node cap, spending budget evenly on good and hopeless branches. Replaced by the control law of §5.
Trajectory storage	Only the final frame is persisted. Fatal here: every decisive observable in the tube record is temporal.
Verdict thresholds	Hard-coded constants on a four-sample estimate. Replaced by explicit, recorded, interval-based promotion rules.
Impossibility reasoning	Currently an <code>if/elif</code> chain that guesses a reason from absent operators. Replaced by an earned claim (§13, defect 8).

2.3 Must be added

The supervisor, the cheap-tier triage, the proximity starvation rule, the watcher, the causal interpreter, and the artefact pipeline — all specified below.

3 The Okuda composition space

The space is defined by an operator vocabulary with typed ports, a stage gate, and a seed. All operators below are already registered and exercised; none is hypothetical.

3.1 Vocabulary

Stage	Operator	Kind	Ports
1 — substrate	<code>seed_mesh_3d</code> / <code>load_mesh_3d</code>	structural	—
	<code>shape_energy_3d</code> (<i>default</i> <i>monolayer</i>)	lateral	<code>a0</code> slot
	<code>reconnect_t1_3d</code>	rewire	—
2 — growth & topology	<code>vesicle_growth</code>	structural	—
	<code>morphogen_growth_3d</code> (<i>conserve_amount</i>)	structural	<code>gate</code> ← morphogen
	<code>divide_3d</code> (<i>hertwig</i> <i>orient_iface</i>)	structural	<code>axis</code> ← morphogen
	<code>extrude</code> (<i>the forcing term</i>)	lateral	<code>site</code> ← morphogen
3 — patterning & feedback	<code>cell_geometry_3d</code>	aggregate	→ geometry
	<code>cell_adjacency</code>	rewire	→ adjacency
	<code>cell_diffuse</code>	lateral	needs adjacency
	<code>cell_react</code> (<i>gray_scott</i> <i>gierer-meinhardt</i> <i>brusselator</i>)	lateral	→ morphogen
	<code>cell_rd_seed</code> (<i>cone</i> <i>tip</i> <i>spot</i> <i>none</i>)	structural	→ morphogen

Two design choices in that table are load-bearing.

Extrusion is an operator node, not a parameter. The whole round-41 finding is “ablate the forcing and the tube collapses”. If the forcing lives inside a configuration file it can only be turned down; if it is a node in the graph, the necessity test **ablates it automatically**, on every composition, as part of the standard protocol. The campaign’s central question thereby becomes a routine measurement rather than a special experiment.

Three reaction kinetics are three implementations of one contract. `cell_react` already resolves to `gray_scott`, `gierer_meinhardt` and `brusselator`. Swapping them is an implementation choice under a single biological contract, which is what the language claims and here is simply true. The search can therefore ask “which kinetics” without changing the composition’s identity, and ask “with or without reaction at all” as a genuine structural edit.

3.2 Seed and legal edits

The seed is Stage 1 only: a mesh, a shape energy, and reconnection — a relaxing vesicle that does nothing. Every subsequent mechanism is reached by one legal edit. Crucially the search visits the degenerate corners on the way (growth with no patterning, patterning with no growth, mechanics alone), and those are where impossibility results come from.

3.3 Phenotype and the metric bank

Quantity	Reads	Answers
<code>aspect</code> , <code>tube_len</code> , <code>n_tubes</code>	shape	is there a tube?
<code>Q</code> (quasi-static)	shape after driver-off relaxation	is it grown or forced?
<code>p_tube/p_body</code>	differentiable stress field	is load shared?
<code>hollow</code> (tube-aware), <code>area_CV</code>	mesh integrity, uniformity	is the wall real?
<code>red_frac</code> , <code>tip_act</code>	activator confinement, activator–radius correlation	does patterning ride the tip?
<code>n_cells</code> , <code>shape_index</code>	proliferation, cell rounding	is it well-formed?

Each is computed *per frame*, not only at the end — see defect 7.

4 Why searches like ours go into rabbit holes

Every defect in §13 is fixable in days. The harder problem is that a search can be free of bugs and still waste months. Our own record supplies the diagnosis, and the pathology has five separable components — worth separating because each needs a different countermeasure.

1. **Depth-first drift.** Three separate rounds tested, in substance, “the same thing but more”: run longer, activate a bigger patch, make the wall thicker. All three were falsified. Each was a locally reasonable next step from the configuration in front of us. None was a new hypothesis. A search that always expands its current best node does this forever.
2. **No terminal-state rule.** Nothing ever said “this branch is finished”. Rounds continued because there was always another knob. Stopping was a human losing patience, which is not a criterion.
3. **Near-duplicate proliferation.** Thirty rounds explored perhaps four genuinely distinct ideas. Without a distance between hypotheses, twenty variants of one idea consume twenty times the budget of one idea — and their collective failure is then misread as strong evidence against the idea, rather than as one observation repeated.
4. **The eye and the number disagreeing, undetected.** Two of the most confident intermediate conclusions were measurement artefacts: a “97% of cells inverted” result produced by pairing each frame’s coordinates with a different frame’s connectivity, and a “101 chemical spots” count where the eye saw three, because the threshold sat inside the noise. Both were

caught by looking at a picture. Neither by a number. In a supervised session a human notices; in an unattended campaign nobody does.


5. **The human as sole goal-holder.** An agent elaborates whatever is in front of it and cannot detect its own drift. Every course correction on record — “we are meant to be testing the growth-driven mechanism, not tuning this extrusion” — came from outside the loop.

The rest of this document is, in effect, five countermeasures.

5 The control law

This is the core of the specification. The governing idea is that budget must be *removed* from unproductive branches automatically, because no one will be watching.

5.1 One round



Stage	Keeps	Rule
Propose	$B = 24$	One legal edit each, drawn across <i>distinct</i> clusters, not all from the current best.
Critic	—	Free rejection: ill-typed edit, unmet precondition, structural hash already evaluated. No simulation.
Tier 1	8	Small mesh, short horizon, seconds each. A gate, never evidence.
Tier 2	8	Full run on the partition, full metric bank, per-frame.
Rank	—	By measured metrics. Pairwise with a Bradley–Terry–Luce aggregation where a scalar order is ill-defined; never win/loss tallies, which carry position bias.
Truncate	3	Losers are dropped, not refined. They stay in the archive as evaluated evidence and are never re-expanded.
Starve	—	Within-cluster competition; a cluster’s members compete against each other, so a family of near-identical ideas exhausts one budget, not twenty.
Freeze	—	A cluster with $> K = 6$ evaluations and no improvement in its best score is frozen; its budget is reallocated the same round.

5.2 Termination and escalation

The supervisor owns the campaign configuration — objective, success criteria, stopping rule, budget — and workers cannot amend it. Each round it computes:

```

if objective_met(ledger):      terminate("objective")
if all_clusters_frozen():      escalate()      # widen stage gate, or request an operator
if dry_rounds >= 2:           escalate()      # no new cluster AND no best-score gain
if budget_spent >= budget:     terminate("budget")

```

Escalate is the important branch, and it is what a human did by hand last time. It means: the current space is exhausted, so either open the next stage gate, or record an **operator request** — a mechanism the proposer wanted and could not express — and move on. Escalation is logged and is a deliverable (§11).

5.3 Rules enforced in code, not recommended in prose

- **A change of numbers is never a new hypothesis** — guaranteed by the structural hash.
- **No manual parameter tweaking outside the loop.** If a configuration needs hand-tuning to work, that is evidence of a missing operator, not of a missing afternoon.
- **No silent caps.** Any round that truncates, samples, or drops candidates says so in the log with counts. A silent truncation reads afterwards as complete coverage.
- **Structural limitations are recorded once and never retried.**
- **Per-round compute is bounded.** Depth is traded against basin width inside a fixed budget, and the trade is recorded. The current driver deepens every four rounds with no

ceiling, which silently converts a campaign of many cheap experiments into a few enormous ones.

6 The scientific protocol: hypothesis first

The control law of §5 says how budget is allocated. This section says what makes an allocation an *experiment* rather than a sample. One rule, enforced in code:

A candidate is not run until its prediction has been recorded.

Every candidate carries, before it reaches the cluster: a falsifiable *claim*, the single *metric* that decides it, an explicit *prediction* on that metric, a *rationale*, and the Grounder’s *citations*. A hypothesis without a prediction is refused at construction.

6.1 Two intents, and why the mixture matters

Each batch is allocated between two kinds of edit:

Confirmatory	an edit the current causal map predicts <i>will</i> work. Consolidates the map. Individually low-information; collectively what makes the map usable.
Adversarial	an edit predicted to <i>break</i> the current best explanation. High-variance; this is where surprise lives.

Because the prediction precedes the run, every outcome falls into a 2×2 :

	confirmed	refuted
confirmatory	consolidates (low information)	surprise (high information)
adversarial	surprise (high information)	breaks as expected (low information)

A batch of pure confirmation (100/0) only ever confirms what we already believe: near-zero information. A batch of pure falsification (0/100) only ever probes the improbable, which mostly breaks as expected and never consolidates a working map. The productive operating point is a deliberate minority of attacks on the current understanding — roughly 70/30, and at times 50/50 when the map is young.

6.2 The 70/30 balance is a setpoint, not a quota

This is the part that makes the discipline machine-enforceable. Define the *surprise rate* of a round as the fraction of resolved hypotheses whose outcome contradicted its prediction. Then the Supervisor closes a loop on it:

surprise	diagnosis	correction
< 0.10	drifted to 100/0: confirming what we already believe	push adversarial
≈ 0.30	the productive band	hold 70/30
> 0.50	drifted to 0/100: everything breaks, no map accumulates	push confirmatory

A quota cannot notice that its questions have become trivial; a setpoint on the prediction-error rate can. This is the mechanism by which the loop avoids both failure modes the hand-run campaign fell into — a long stretch of confirmatory tuning where nothing surprised us, followed by a stretch of wild attempts that mostly broke as expected.




Figure 1: The fifteen agents. **Red border** = a genuine language-model call (eight: Proposer, Reflection, Judge, Analyst, Watcher, Interpreter, Evolution, Meta-review). Grey = deliberately deterministic, following Robin’s de-agentification — the Critic is the type guard and so cannot itself be a model, the Referee ranks by measuring, and an LLM Metrologist could talk itself into admitting its own evidence. **Green** = the human: the Engineer patches the substrate from outside the loop, because an agent able to rewrite the simulator can make any hypothesis true.

6.3 The knowledge document

`knowledge.md` is append-only and written in the order the science happened: the hypothesis and its grounding, then its prediction, then what occurred. A reader can audit what was believed *before* the evidence arrived — which is the only way to tell a prediction from a post-hoc rationalisation. Verdicts move between the four classes of §11; a claim’s history is extended, never rewritten. Surprises are flagged, because they are the entries a returning human should read first.

7 The agents

Two systems published together in *Nature* (July 2026) solve much of the orchestration problem for literature-driven hypothesis generation with wet-lab validation: Co-Scientist, a supervisor-coordinated coalition with an Elo tournament, a proximity graph and a meta-review that appends learned patterns to other agents’ prompts; and Robin, which batches and truncates, ranks by an LLM-judged pairwise tournament aggregated with a Bradley–Terry–Luce model, and runs eight independent analysis trajectories reconciled by consensus.

Our setting differs in one decisive respect. **Their experiment is scarce and ours is cheap.** A wet-lab assay costs weeks, so ranking by argument is rational. Our simulation costs seconds and is measured by code. We therefore take their orchestration, diverge on their adjudication, and add one agent neither has.

#	Agent	LLM	What it does	Where it lands
1	Supervisor	—	holds objective, budget, terminal state; escalates	<code>control.py::Supervisor</code>
2	Grounder	part	<code>plexus2 + okuda + corpus</code> ; <code>ground / gate / understand / figure.target</code>	<code>agents/grounder.py</code>
3	Proposer	yes	chooses the edit from a type-filtered menu; commits to a prediction	<code>agents/proposer.py</code>
4	Reflection	yes	peer-reviews the batch <i>before</i> it costs cluster time	<code>agents/llm_agents.py::reflect</code>
5	Critic	—	types, preconditions, duplicates — free rejection	<code>composition_space.py</code> , <code>propose_batch</code>
6	Referee	—	ranks by <i>measuring</i> , gate-admissible metrics only	<code>control.py::score_run</code> , <code>rank_btl</code>
7	Judge	yes	second opinion from the pictures; disagreement is the signal	<code>llm_agents.py::judge_pair</code>
8	Proximity	—	structural distance; near-duplicates starve together	<code>control.py::ProximityIndex</code>
9	Analyst $\times N$	yes	N independent readings $\rightarrow$ consensus; disagreement recorded	<code>llm_agents.py::analyse</code>
10	Watcher	yes	captions the movie (VLM) and veto es promotion	<code>run_one.py::describe</code> , <code>llm_agents.py::watch</code>
11	Interpreter	yes	this spec $\rightarrow$ this phenomenon, by this route	<code>llm_agents.py::interpret</code>
12	Evolution	yes	refines the winner: simplify / synthesise / ground	<code>llm_agents.py::evolve</code>
13	Meta-review	yes	distils patterns into the Proposer’s prompt (capped)	<code>llm_agents.py::meta_review</code>
14	Metrologist	—	certifies the instrument; issues retractions	<code>agents/metrologist.py</code>
15	Engineer	human	patches the substrate, from outside the loop	—

Table 1: The roster in full. Eight genuine language-model calls; five deliberately deterministic; one part-deterministic (the Grounder’s retrieval is reproducible by construction, so a citation can be re-derived); one human. The shared substrate — the CLI wrapper, the file-based working memory (`instruction / memory / append-only analysis / user_input`), the wall-clock budget, forced seeds and the causality rule with its control slot — is `agents/llm.py`, ported from an earlier exploration loop that had already solved these.

Agent	Role	Provenance / divergence
Supervisor	Holds objective, criteria, stopping rule, budget. Computes round statistics; decides allocation, freezing, escalation, termination. Never does the work.	Co-Scientist’s Supervisor, whose statistics likewise determine whether a terminal state is reached.
Grounder	Holds <code>plexus2.tex</code> (the language contract), <code>okuda.pdf</code> (the reference model) and <code>okuda_corpus.md + papers/</code> (the literature). Answers mechanism questions and attaches citations.	Robin’s literature agents. Co-Scientist’s ablation showed that giving the reviewer search suppressed plausible-but-false hypotheses.
Proposer	Generates a batch of one-edit candidates across distinct clusters.	Co-Scientist’s generator. <i>Divergence</i> : our space is typed and executable, so illegal proposals are impossible by construction rather than rejected on review.
Critic	Free pre-compute triage: types, preconditions, duplicates.	Co-Scientist’s reflection, reduced to what can be decided without simulating.
Referee	Ranks a batch.	Robin’s tournament shape and BTL aggregation. Divergence: the judge is the metric bank, not a debate.
Proximity	Clusters compositions by structural distance for within-cluster competition.	Co-Scientist’s proximity graph. Ours is <i>exact</i> for identity (the hash) and approximate only for similarity.
Analyst $\times N$	N independent readings of the same trajectory, reconciled by consensus; disagreement is itself recorded	Robin’s eight parallel analyses. Aimed squarely at our two measurement artefacts.

The Grounder, specifically. It is worth being precise about this agent because it is the one whose absence was most expensive last time. The Grounder holds three bodies of text — `paper/plexus2.tex` (what the language permits), `papers/okuda.pdf` (what the reference model actually does), and `papers/okuda.corpus.md` together with the vendored `papers/` corpus (what the field has already established) — and it is called at exactly three points:

1. when the **Proposer** needs a mechanism grounded before proposing it;
2. when the **Supervisor** gates a hypothesis, before Discovery is allowed to propose;
3. when an **operator request** is filed — does the literature already name this mechanism, and under what conditions does it hold?

This is the agent that produced the Gierer–Meinhardt unlock by hand last time. The decisive question — *which reaction–diffusion regime does Okuda actually use?* — was answered by reading the paper, not by sweeping parameters, and it redirected the entire campaign. Making that a scheduled call rather than a lucky afternoon is the point.

Two further borrowings worth naming. Whatever proposes must not be what judges — Robin used different models for generation and adjudication, and the separation is free. And Robin found their orchestrator “almost always called tools in the same order”, so they replaced the agentic controller with a deterministic notebook. We follow that: **the control law of §5 is a deterministic script**, and language models are called only where judgement is genuinely required — proposing, watching, explaining.

One thing neither paper needs. Their measurement is an external assay. Ours is code written by the same system being judged, so “the metric lied” is a failure mode unique to us, and we have suffered it twice. Hence the instrument gate of §13.

8 Knowledge certification, and who may touch the substrate

If the foundations are wrong, so is the knowledge.

Every other agent in §7 produces or judges *evidence*. None of them is responsible for the thing evidence is measured *with*. That gap is not hypothetical: it is how the defect below was found — by a person reading code, not by any designed role.

8.1 A worked example of why the role must exist

`divide3d` counts `min_cycle` and `max_cycle` in *division-calls*, and `max_div / max_div_frac` are *per-call* throttles. The archived configurations passed **every**: 2, which the engine gated *and* the operator’s own private counter gated — so the operator body executed once every **four** frames. Correcting the clock therefore multiplied the wall-clock meaning of every per-call quantity by four, and the same recipe that had produced a tube of aspect 7.5 produced one of aspect 3.2.

Two further consequences, each of which a casual fix would have missed:

- The hand-tuned values had been adopted as *vocabulary defaults*, on the reasoning that “a default composition should start where our evidence is”. After the clock fix that reasoning inverted: those defaults now encode a θ tuned for the wrong clock, so *every* composition the loop generates would have started biased — silently.

- Rescaling the obvious knob was not enough. Since $\text{cap} = \max(\text{max_div}, \text{max_div_frac} \cdot n_F)$, the absolute floor *dominates* at realistic cell counts, so rescaling only `max_div_frac` was **entirely masked**. Both had to move.

Because the factor is exact, the correction is analytic rather than a re-tuning sweep: rescale the per-call quantities and the fix becomes *behaviour-preserving by construction*, verified identical to the archived budget at every cell count. Any later change in phenotype is then attributable to the change that caused it, rather than to the clock.

8.2 The Metrologist

The Metrologist certifies that evidence may be admitted at all. It owns the validation ladder, the instrument gate, and — the part no other role covers — *substrate semantics*: units, clocks, and the distinction between per-call and per-frame quantities. Its distinguishing power is that it may act *backwards*:

When a foundation defect is found, the Metrologist issues a **retraction**: every claim whose evidence depends on the defective semantics is moved back to *Open*, with the defect and its scope recorded.

In an append-only ledger a retraction is a new record, never an edit. The history must show what was believed, on what evidence, and why that evidence was later withdrawn — otherwise a reader cannot distinguish a claim that survived scrutiny from one that was quietly repaired.

8.3 The boundary: who may modify the substrate

No campaign agent may modify the substrate. An agent that can rewrite the simulator can make any hypothesis true.

The separation is therefore structural, not procedural:

Metrologist	<i>detects and quarantines.</i> Halts admission of evidence, marks affected claims Open-pending-recheck, and states the scope of the defect. It does not write code.
Engineer	<i>proposes a patch,</i> into the substrate — a <i>different artefact</i> from the campaign’s evidence. It cannot admit its own fix.
Supervisor	<i>gates resumption.</i> Evidence admission restarts only after the full ladder passes on the patched substrate, and the re-anchoring is recorded explicitly rather than absorbed.
Human	approves substrate changes. The campaign runs unattended; the thing it measures with does not change unattended.

The underlying principle is the one an experimentalist would recognise: *the instrument must not be adjustable by the experiment it is measuring*. Our situation is sharper than either published system’s, because their measurement is an external wet-lab assay and ours is code written by the same project that is being judged.

9 Which agents are language models, and which must not be

Co-Scientist runs six specialised LLM agents under a mechanical supervisor: Generation, Reflection (peer review *with search*), Ranking (Elo by simulated scientific debate), Proximity, Evolution (refines the top-ranked) and Meta-review (distils patterns and appends them to the other agents’

prompts — “feedback propagation without back-propagation”). Robin uses Crow and Falcon for literature and Finch for analysis, ranks with an LLM judge over pairwise comparisons aggregated by Bradley–Terry–Luce, and runs *eight independent* Finch trajectories on the same data because the analysis “can vary between runs, even when given identical prompts and data”.

We follow that, with two deliberate departures and one addition.

Agent	LLM?	Why
Proposer	yes	Chooses which mechanism edit is worth a GPU-hour and commits to a prediction. An enumerator cannot do this: if the prediction follows mechanically from the edit’s sign, the surprise rate measures nothing.
Reflection	yes	Peer-reviews the batch before it costs cluster time. A type-correct proposal can still be a bad experiment.
Analyst $\times N$	yes	N independent readings of one run, reconciled by consensus; disagreement is recorded as a calibration signal.
Watcher	yes (VLM)	Reads the movie and <i>veto</i> es promotion when the picture contradicts the claim.
Interpreter	yes	Writes the causal account — this spec, this phenomenon, by this route.
Evolution	yes	Refines the winner (simplify / synthesise / ground).
Meta-review	yes	Distils recurring patterns into the Proposer’s instructions.
Judge	yes	A <i>second opinion</i> on ranking, from the pictures.
Supervisor	no	Arithmetic on statistics. An LLM here is drift, not judgement.
Critic	no	The type and precondition guard. It is the thing that stops a language model proposing something ill-typed; it cannot itself be one.
Referee	no	Ranks by <i>measuring</i> .
Proximity	no	Structural distance — exact, not approximate.
Metrologist	no	Certifies the instrument. An LLM here could talk itself into admitting its own evidence.
Engineer	human	The substrate boundary.

Departure 1: ranking is measured *and* judged. Both papers rank by argument because a wet-lab assay is scarce. Ours is a simulation scored by code, so an opinion about available ground truth is a step backwards — *except* that the instrument gate has now measured a metric scoring 9.30 on a small bud. So the metric bank ranks first, an LLM judge ranks independently from the rendered movie, and **their disagreement is the signal**: it is the eye/number divergence detector. Neither paper needs this, because their measurement is external to their system; ours is code we wrote.

Departure 2: pre-registration. Neither system records what was *believed before* the evidence arrived, so a prediction cannot be distinguished from a post-hoc rationalisation. Ours refuses to run a candidate whose prediction is not already recorded (§6).

Addition: the type system is not an agent. Illegal edits, unmet preconditions and dangling slots never reach the Proposer at all — it chooses from a menu the type system has already filtered. The language decides what is *possible*; the agent decides what is *worth doing*. Neither can do the other’s job.

10 Running for weeks: where this breaks

The campaign is meant to run unattended for weeks. That is a different engineering problem from running one round, and the arithmetic should be done before the fact rather than discovered.

One run is roughly 250 s of engine, 60 s of rendering, 60 s of captioning and a mechanics pass: call it **7 minutes**. Eight-way on the partition gives ~ 68 runs per hour. A round is a batch of eight plus a proposer call, so ~ 20 minutes, or ~ 70 rounds per day. Two weeks is then of order 10^3 rounds and 10^4 runs.

	Bottleneck	Consequence, and what it forces
1	The composition space is too small.	Twelve operators, ~ 19 legal edits from the seed. Distinct <i>clusters</i> number in the tens. At eight runs a round the search exhausts itself and freezes every cluster within days . A multi-week campaign is therefore <i>not</i> mechanism search — it is dominated by parameter fitting and by operator requests . Escalation is the main path, not an edge case, and it is the least built part of the system.
2	VLM cost	A 23 GB model loaded once per job. Caption per <i>wave</i> instead: one load, eight captions. Still always-caption, never deferred to an end pass.
3	Storage	30–50 MB per trajectory $\times 10^4$ runs ≈ 400 GB. Retention policy: full trajectories for kept and promoted runs, per-frame metric tables (kilobytes) for everything else — so the archive stays re-scorable exactly where re-scoring matters.
4	The proposer call blocks	Eight GPUs idle for minutes every round. Propose round $N+1$ while N runs.
5	Knowledge growth is unmeasured	The point of the campaign is that knowledge accumulates, and there was no quantity tracking it. See below.
6	The videos do not accumulate	Each run has its own movie; nothing assembles them into something you can watch. See below.
7	Prompt growth	Meta-review appends forever. The distilled section is capped and rewritten in place, so old patterns are dropped rather than accreted.

10.1 Knowledge growth, as a measured quantity

If the campaign’s product is knowledge, then its rate is the thing to watch, and it should be plotted rather than asserted. Per round we record: distinct clusters discovered and frozen; counts moving between Established / Refuted / Structural / Open; the **surprise rate** (§6); operator requests filed; and the fraction of runs admitted as evidence at all. A campaign whose surprise rate collapses toward zero has stopped learning even while it is still burning GPU-hours — and that, not a wall-clock limit, is the honest signal to escalate or stop.

10.2 The video battery

Every run already produces a movie, a strip, a caption and a mechanics panel. What was missing is that they never assembled. Three artefacts accumulate across the campaign:

- a **round montage** — the batch’s final frames side by side, so a round is legible at a glance;
- a **progress reel** — the best run of every round, concatenated in order, so the campaign can be *watched* learning;

- a **phase-diagram panel** — the four Okuda cells (§1) as they currently stand, which is simultaneously the scientific target and the talk figure.

The reel is not decoration. The single most expensive failure in this project’s history was a number disagreeing with a picture, and the cheapest defence has repeatedly been to look.

11 What the campaign produces

Five artefacts, each a deliverable in its own right.

11.1 The ledger

Every claim sorted into **Established**, **Refuted**, **Structural limitation** or **Open**, with its evidence, sample count, confidence interval and the threshold set that produced the verdict. A mechanism reaches Established only if it answers what a biology paper would ask: remove it — does the phenotype collapse? Is a minimal recipe containing it sufficient? Does it survive repeats and a range of settings?

11.2 Causal descriptions — “this spec gives this phenomenon”

For every surviving composition the Interpreter writes a structured account, attached to the record as a versioned analysis:

```
{ composition: <structural hash + operator list + typed connections>,
  phenotype:   <class, with the metric vector that assigns it>,
  route:       "growth gated by the activator raises target volume locally;
               oriented division builds a wall along the bud axis; T1 lets
               the wall cells slide, so the tube extends instead of bulging",
  necessary:   [ops whose ablation collapses the phenotype],
  sufficient:  <the minimal sub-composition that still produces it>,
  falsified:   [alternatives tried and refuted, with run ids],
  scope:       <substrate, parameter ranges, seeds>,
  watcher:    <caption of the movie>, agreement: <caption vs metrics> }
```

This is what makes the operator library reusable rather than merely catalogued, and it is what a reader of the eventual paper actually wants. It is generated per composition, continuously, not written up afterwards.

11.3 The operator backlog — atlas growth

Every escalation records an **operator request**: a mechanism the proposer wanted and the language could not express. Each is a candidate new operator and each is a talk bullet. Our expectation, from the one simulator decomposed in full, is that mechanics *saturate* against the existing algebra while topology and numerics contribute genuinely new contracts — and that saturation is itself the hypothesis under test. If genuinely new families keep appearing, the language is incomplete; if they stop, the compactness claim is supported.

11.4 Movies

Every promoted composition is rendered to the standard of the public gallery, captioned by the watcher, and archived beside its record. Captioning is on by default and is never disabled: the caption is simultaneously documentation, a semantic regression test, and the watcher’s input. These

are the primary evidence for a talk, and the loop produces them rather than a person producing them afterwards.

11.5 One differentiable figure

See §12.

12 Where the differentiable version enters

Mechanism search and parameter fitting must not run simultaneously, or neither conclusion is interpretable. The fit is *earned*, and the trigger is explicit:

A composition that is **Established** but whose residual on a chosen objective persists *across the whole parameter basin* is handed to the differentiable fit. A residual that survives every setting is evidence about a mechanism; a residual that disappears at some setting was never a mechanism question.

The objective is local in space and time. A loss spanning every cell over hundreds of frames unrolls, under reverse-mode differentiation, into a graph whose memory scales as cells $\times$ steps and rapidly becomes intractable. But the evidence constraining any single operator is usually confined to a small, transient region — the advancing front, the neck, the moment a wall forms. So the loss is placed exactly there: a region of interest over a few tens of steps. The mechanism is global and shared; the objective is not.

The natural objective here is the stress field. The tool that computes per-cell force, pressure and cortical tension is already torch-differentiable, and the campaign’s discriminator Q is mechanical. “Minimise tube pressure while preserving aspect” is therefore both the biological question (is a low-stress tube reachable?) and a well-posed differentiable objective.

Deliverable. One figure: residual descent, with identifiable parameters distinguished from degenerate ones. A degeneracy found is a result — it says two mechanisms are indistinguishable under the current observable, which is a measurement question, and it goes back to the observable bank.

13 Pre-flight: eleven defects that must be closed first

The campaign will emit scientific claims, including claims that something is *impossible*. A false impossibility claim is worse than no claim, because it closes a question that should stay open. Defects 1–6 come from an adversarial audit of the simulation operators; 7–11 were found while writing this document, by reading the loop’s own source with the same suspicion.

1. Two operators kept a private clock. They spaced themselves out while the engine *also* spaced them out; the schedules multiplied, so “every 2 steps” fired every 4. Division across every archived run happened at a fraction of the advertised rate. *Fix: the engine owns the clock; baselines re-measured and the re-anchoring stated, not absorbed.*

2. Adding a signalling step silently changed the time-step. The run script chose the time-step from which operators were present — recipes with a chemical signal ran fifty times finer. Since the campaign’s most important edit is adding or removing signalling, every conclusion about whether signalling matters would have been confounded. *Fix: one time-step for the whole search; stability handled inside the operator that needs it.*

3. A silent mismatch produced a result we believed for days. Positions and connectivity were recorded on different strides; the analysis paired coordinates with another moment’s connectivity and reported cells as inverted. Every physical cure failed — in hindsight, the clue. *Fix: equal length asserted, not clamped.*

4. An operator that does nothing still returns numbers. If a prerequisite is missing, the operator silently no-ops, the run finishes, and a score is produced. Harmless under hand-written recipes; under automatic search it is the most dangerous bug possible — the loop records “this mechanism cannot make tubes” when the mechanism never ran. *Fix: preconditions declared and checked; every operator reports whether it acted; a scheduled operator that never acted is an error, not a result.*

5. Some operators advertise capabilities they lack. One is labelled differentiable and discards gradients; another claims three dimensions while its validity check is two-dimensional. Labels that lie are worse than absent labels, because the search trusts them. *Fix: corrected and mechanically certified.*

6. Measurements can pass while the picture is wrong. A lumpy blob scored acceptably on every number we had, because the scores measured how far things stick out and a blob sticks out. *Mitigation: the watcher, plus the instrument gate below.*

7. The archive stores only the final frame. Each run computes several recorded frames and persists only the last, so every re-scoring is restricted to end-state observables. This contradicts the entire record: the decisive evidence was temporal — a protrusion that peaks mid-run then degrades, a pattern that floods only once growth starts, divisions arriving in synchronised waves. Worse, the artefact that cost us days is by construction a time-series defect, and an archive of final frames cannot be re-examined for it. *Fix before any long run: persist the full trajectory and a per-frame metric table. Storage is cheap; re-simulating weeks of cluster time to recover a time course we already computed and discarded is not.*

8. Impossibility is asserted by a hard-coded narrative. The strongest claim the system makes is currently produced by a low score plus an `if/elif` chain guessing a reason from absent operators. The reason is a template, not evidence; combined with defect 4 this is machinery for manufacturing confident false impossibility. *Fix: an impossibility claim requires (i) every scheduled operator verified to have acted, (ii) a minimum sample count with an explicit interval, (iii) a recorded, bounded falsification attempt — a deliberate search for any setting that does produce the phenotype — and (iv) a stated scope. Absent any of the four, the verdict is Open.*

9. Four samples cannot support a robustness claim. The default basin is two seeds by two parameter draws, so a rate is one of $\{0, \frac{1}{4}, \frac{1}{2}, \frac{3}{4}, 1\}$, and promotion thresholds are hard-coded constants — Established at ≥ 0.6 (three of four), structural below 0.2 (zero of four), necessary

if ablation drops below 0.5. *Fix: a minimum sample count before promotion; a Wilson interval rather than a point rate; thresholds as named, recorded campaign parameters so a later re-analysis can vary them.*

10. The search accelerator is dead code in the long-run driver. Loop I accepts an optional surrogate that ranks candidate edits so the frontier expands best-first; the overnight driver never passes it. Every unattended round to date has been uniform breadth-first — the exact opposite of §5.

11. Per-round cost grows without bound. The driver deepens the search every four rounds and densifies the basin every twelve, with no ceiling.

The instrument gate

Before the first round, the metric bank must correctly separate a set of already-archived runs we have labelled by eye — a clean tube, a capped lobe, a flooded shell, a uniform sphere, a cauliflower. If the instrument cannot reproduce our own judgement on known cases, it is not ready to judge unknown ones. This is not a formality: three metrics in the existing bank have been caught lying, one of which moved a headline number by thirty per cent once corrected.

The validation ladder

One command, each level independently pass/fail, non-zero exit on any failure.

Level	Checks	Error class caught
Contract	every operator declares what it is and does	labels that lie
Unit	each operator reproduces its previous output	the refactor changed physics
Trajectory	a run reproduces its recorded history	schedule or ordering drift
Invariants	closed surface, no tangled or bow-tie faces	silent geometric corruption
Recording	positions and connectivity are equal length	the phantom-result bug
Metrics	scores match archived values	the analysis path drifted
Gradients	derivatives flow and are finite	the fit will fail later
Coverage	every operator exercised by some configuration	untested code presumed working
Description	the watcher’s caption matches expectation	the numbers pass, the picture is wrong

One rung deserves comment. A face must not be self-intersecting, and we had checked validity by requiring positive area. *A bow-tie has positive area.* The check passed, the mesh was tangled, and a quantitative result about cell rearrangement was wrong — caught by watching a movie, not by any number. The lesson generalises well beyond meshes: a numerical invariant is not the geometric one you meant.

14 Operating it

The partition is cheap enough that wall-clock is not the binding constraint — attention is. The operating rules follow from that.

- **Read the ledger, not the runs.** The weekly unit of human attention is the ledger delta — what moved between the four classes — plus a montage of the week’s best and the current operator backlog. Individual runs are opened only when the ledger points at one.

- **Verify against the world, never against a submission’s reply.** A submission that reports failure may have succeeded, and one that appears to time out may already be queued; we have created duplicate jobs this way. The queue is the only ground truth. Generalised: an action’s *reported* outcome is a hint; the world’s state is the fact.
- **Every round is restartable.** Evidence is persisted as produced, exceptions caught per round, run identity a content hash — so a crash resumes without duplicating or losing evidence.
- **The configuration is written before the first round,** not discovered during it.

15 Status and the switch-on gate

The loop has run on gland branching and produced, unattended, a sequence that included one conclusion overturned by a better measurement — the strongest evidence that the evidence/interpretation separation is worth its cost. The Okuda operators exist, have produced a tube, and have produced the more interesting negative result that our tube is forced where the published one is grown. That is the question the campaign inherits, and §1 turns it into a measurable one.

The campaign switches on when four things are true:

1. the eleven defects are closed, each demonstrated by a test that fails loudly;
2. the archive stores trajectories and per-frame metrics, not final frames;
3. the metric bank has reproduced our own eye-labels on already-archived runs;
4. the supervisor’s campaign configuration — objective, criteria, stopping rule, budget — is written down.

That gate is the point of this document. The system is now capable of running unattended for a long time and producing confident, well-formatted, entirely wrong conclusions. The defence is not more compute. It is a small number of checks that fail loudly, an archive that keeps raw evidence separate from its interpretation, a search that drops its losers instead of lovingly refining them, a supervisor that holds the goal nobody else can be trusted to hold — and the willingness to keep looking at the pictures.

16 First live results (2026-07-30, overnight)

Three things were exercised for the first time: the LLM Proposer, the composition mode end to end, and the loop pointed at a *second* system (the Turing × vertex recipe behind the site’s front-page movies). Each produced defects that only running could find.

16.1 The composition mode could never have completed a round

`proposer.propose()` *rejects* any proposal whose slot 0 is not `intent: "control"; hypothesis.INTENTS` did not contain `"control"`, so `Hypothesis.__post_init__` raised on it. Two halves of the same protocol disagreed about whether a control exists and both enforced their view with a hard error. A control is now a first-class intent, excluded from the surprise rate and the mix — it asserts nothing about an edit — but a control that *fails* its own prediction raises `BASELINE DRIFT`, which is louder than a surprise: every other slot is measured as a difference from it.

16.2 The surprise rate was computed by a regex that could only bias it downwards

The campaign’s one control signal was a first-match regular expression with three defects: an *unparseable* prediction returned true and was recorded **confirmed**; the parser was metric-blind, applying the first comparison it found to **protr_peak** whatever the clause named; and **REFUTED** if ... was read as the assertion. All three push the rate toward “nothing was learned”.

The measured cost: the **vcap** round’s five points were posed before the parameter-hypothesis fix with the placeholder **"unknown -- sensitivity sweep"**, which parsed to nothing and so resolved **confirmed**. The persisted surprise rate for the campaign’s most informative round was **0.00**. The 0.40 everyone quoted had been computed by hand in a narrative and never reached the ledger. **predict.py** now parses every clause with its metric and returns **inconclusive** rather than guessing; an append-only **amend()** backfilled the record, and the ledger computes 0.40 itself. The same guard later refused to score 32 runs of the Turing $\times$ vertex study because **protr** was not a known metric — declining 32 times instead of inventing 32 confirmations.

16.3 The supervisor did not survive a restart

round, **spent**, **dry**, **best**, **prox** and **stage_gate** were re-initialised every process; **supervisor.jsonl** was append-only *output*, never read back. For a campaign specified to run for weeks across restarts this disabled every stateful control at once: every round is “round 1” (five duplicate hypothesis ids were already in the ledger), the budget can never be spent so **terminal()** never fires, dry rounds never accumulate so **escalation was dead at the trigger**, and no proximity cluster is ever frozen. State is now checkpointed atomically, and a corrupt checkpoint refuses to start rather than silently resetting the counter.

16.4 The escalation path, and its first entry

terminal() has always been able to return **ESCALATE**: and nothing ever read it, so escalation could only be reached by a crash-like condition — never by a campaign exhausting its reachable space, which over weeks is the normal ending. The path is now three actions tried in order (open a stage gate; file an operator request; declare the region exhausted) and it has produced **OR001**: quasi-static force-balance relaxation to a *residual* rather than a fixed step count, plus 3D IH/HI reconnection, with an acceptance test that today’s best setting fails by roughly $4\times$ on both integrity metrics.

16.5 Turing $\times$ vertex: a structural limitation, measured

The two front-page movies are the *same simulation*: their archived diagnostics are identical because **a_sw = 50** leaves the Hill switch unfired, so the reaction–diffusion is pure colouring. Running the loop on the settings found something sharper — with **rho = 1.0** every cell grows every tick, and the **conserve_amount** dilution then drives the activator to 2×10^{-22} by frame 400. The coral pattern on the site’s front page is a colour map stretched over an extinct field.

With the field alive, 32 runs across five knobs give a clean boundary:

ρ	1.0	0.5	0.2	0.05	0.0
corr(act, radius)	+0.29	+0.26	+0.28	+0.63	+0.53
protr	1.03	1.05	1.06	1.17	32.7
hollow_frac	0.000	0.006	0.008	0.197	0.839

Eighteen runs keep an intact mesh (`hollow_frac` < 0.05) and reach at most `protr` 1.064; all nine runs with `corr` > 0.4 have `hollow_frac` $\geq$ 0.197; **none achieves both**. On this backend activator-driven shaping and monolayer integrity are mutually exclusive. That is a structural limitation, it is first-class, and it reaches by measurement the conclusion the previous campaign reached independently on a different backend: the gap is representational, not parametric.

16.6 What the defences caught, and what they cost me

On the first composition round the Watcher vetoed the two *highest*-scoring runs — `protr_peak` 4.03 and 3.20, one of them the control that three Analysts had called a tube — because the captions show no tube. The Critic separately refused a third slot as not-evidence (reservoir saturated at 15 002 cells). Two of six slots produced no evidence and both said so loudly.

Twice in one night I proposed a mechanism and the measurement refuted it. I argued that the per-cell volume cap explained the flat shape; the first test refuted it, but that test ran against the extinct field, so its refutation was worthless — *a refutation obtained under a broken instrument is not evidence, and I nearly filed it as one*. I then argued the extinction reached the main campaign, submitted the comparison, and measured `act_max_final` = 0.995 with the dilution on: refuted within the hour. The dilution tracks the volume growth ratio, so it only bites while cells grow; at $\rho = 0$ growth saturates and the dilution stops. I had the sign of the feedback backwards.

Both are recorded as retractions rather than quietly edited. The protocol’s value is not that it prevents wrong ideas — it is that a wrong idea costs two cluster runs and an hour instead of a month.

17 External review — 2026-07-31, 07:58 EDT

Commissioned against the goal as stated: “understand the Okuda system / reproduce the paper figures mechanically in Plerus2.” Six reviewers covered fidelity, metrology, the control law, the agents, code correctness and strategy; every non-minor claim was attacked by an adversarial verifier and the load-bearing items re-checked by hand against code and data.

Verdict. The instrument is better built than it is aimed, and three load-bearing measurements are broken in ways that make the current record uninterpretable. The discipline — pre-registered predictions, job-ID tracking, the inert-operator ledger, a parser that refuses to guess — is above normal practice and must survive what follows. But the coded objective (lever-map coverage) is not the stated goal, and the deliverable of §1 is not reachable through the coordinates in use.

1. Q measures the initial condition. `quasi_static_Q` (`run_one.py:271--303`) drops growth and driver, then calls `S.load(spec_q.yaml)` — a *fresh* simulation — so it relaxes the seed sphere for 60 frames instead of continuing from the end state. Across `log/okuda/*/diag.json`: $Q = 1.014$ in **14 of 16** runs, whose `protr_final` spans 1.02–2.81. Q has weight 1.0 in `score_run` and gates `meets_success` at $Q \geq 2.0$: the success criterion is unreachable and the objective reduces to `0.5*protr_peak` plus a constant. *Every forced-versus-grown verdict in the campaign is narrative over a null instrument.*

2. The 32-run structural-limitation study was rail-limited. All 32 runs end at `cells_end = 1778` exactly while mid-run counts differ (679–683). The vertex reservoir is `int(Nv*12) = 3552`; for a trivalent spherical mesh $V = 2F - 4$, so F is capped at $(3552+4)/2 = \mathbf{1778}$. Arithmetic, not coincidence: *every run terminated against the vertex buffer* — the `P2_BUFFER_SATURATED` condition the main campaign calls not-evidence and `turing_vertex_study.py` never checks. Reproducing the archive’s 1778 thus validates the clock migration only in the weak sense that both hit the same ceiling. OR001 may well be right — Okuda is explicitly quasi-static — but it is *not yet measured*. Re-run at a 10× reservoir before inheriting the claim.

3. The `ta_*` family measures radius from the world origin. `tube_analysis.py:22` returns `norm(cen,axis=1)` with no centroid subtraction, and the cluster directions inherit it, while `run_one.py:108` `protr_of` *does* subtract `pos.mean(0)`. `protr` and `ta_protr` are therefore different quantities — the root cause of the “my metric vs. their metric” rule — and the admitted `ta_n_tubes_final` inherits it.

The goal is unreachable through the named coordinates. `chi` (`composition_space.py:180`) multiplies `d.a` and `d.h` *equally*: it moves pattern wavelength but is not Okuda’s χ , and its range (1,10) sits below the 30–900 `calib_chi.py` needed for his spot count. `gamma` is emitted only in the `else` branch of `_emit_react` (`translate.py:136--147`) — Brusselator only — so under the Gierer–Meinhardt kinetics the campaign says Okuda uses, *the γ knob does not exist*; `rd_rate` $\in (0.2, 3.0)$ spans 1.2 of the required 4 decades. Both operators are stage-3, outside the move set, and `instruction.md:36` forbids parameter edits as hypotheses. No admitted metric separates undulation from tubulation from branching, and `ta_tube_diam_final` — the thin/thick axis, the one observable χ controls — never reached the gate. **Calibrate to Okuda’s observables (spot count, tube diameter), not to his parameter values.**

Control-law defects that fire on the next round. `escalation.decide` checks `frozen_all` *before* `backlog.open_requests()` (`escalation.py:180--183`), so with OR001 open and the single cluster frozen it returns `request_operator` every round forever; `exhausted` is unreachable. `round.py:344` saves only `kept` and a Watcher veto scores $-\infty$, so `frontier.json` now holds exactly the two knockouts (`-extrude`, `-morphogen.growth_3d`; `protr` 1.39, 1.03) and has *discarded* the parent 5e3159, the only composition that produced a protrusion. The veto is also circular: `round.py:291` passes `analyst_consensus` — a label derived from the caption the Watcher reads — where the hypothesis’s claim belongs, so it detects analyst-caption drift, not eye/number divergence. `run_id` passes `params={}` (`run_one.py:197`): all five vcap runs share one archive record.

What is solid. The Analysts (the only role seeing pixels and raw numbers); the cluster driver (job IDs, `bjobs -a`, straggler kill); `predict.score`, `critic.py`, `instrument.py`; the monolayer energy, Okuda Eq. (3) verbatim with bending stiffness $\sim \kappa_s h^2$ emerging rather than added; the Grounder’s figure table; and the pre-registration/retraction discipline — the campaign’s genuine contribution.

Recommended sequence. Fix or remove Q ; fix the `ta_*` centroid and re-run the instrument gate; re-run the 32-run study at a 10× reservoir (one wave — it gates everything downstream); implement OR001 as Okuda’s *published* method rather than a discovery, with the absent deformation→patterning arrow (s_{jk}/v_j weighting, today a degree-normalised Laplacian); author the four discriminating metrics; then run the designed (χ, γ) grid — six jobs, one wave. The loop re-enters afterwards, where it is strong: ablation and necessity testing against a known-good reference. Running unattended before Q and the centroid are fixed writes uninterpretable evidence into a permanent ledger.

18 Plan of record — 2026-07-31, 08:05 EDT

Agreed with C.A. after §17. Two deliverables, in order: (1) *the method*, (2) *the Okuda figure that demonstrates the method*. The figure is not a side quest — it is the evidence that the loop works. Calibration is to Okuda’s observables (*spot count*, *tube diameter*), but the settings follow his Table 1 as closely as the substrate allows.

18.1 Three diagnoses that set the order of work

The phantom guard is not installed where the phantom lives. `check_alignment` (`run_one.py:70--79`) raises when `len(positions) ≠ len(topology)` — the guard written after the phantom “97% hollow” result. `turing_vertex_study.py` uses the exact pattern that guard forbids, `hist[min(t, len(hist)-1)]`, at lines 218, 224 and **254 (the render path)**, and never calls it. It may be a no-op today — the study pins `topo_snapshot_3d` `every=1` with `record_cap = frames+2` — but nothing verifies it, and that is exactly how the 97% result happened. *Every hollow_frac behind OR001 came through the unguarded clamp*. One assertion decides whether “shaping destroys the mesh” is physics or an artefact.

What Q actually computes. Q is specified as: take the end state, switch growth and forcing off, continue mechanics, measure the surviving protrusion. The code deletes those operators from the spec and then calls `S.load(spec_q.yaml)` — constructing a *new simulation from the seed sphere*. It therefore relaxes a fresh sphere for 60 frames. Hence $Q = 1.014$ in 14 of 16 runs whose true `protr_final` spans 1.02–2.81. Q needs a *checkpoint*, not a re-spec.

What `ta_*` actually computes. `tube_analysis.py:22` takes each cell’s radius as `norm(centroid)` — distance from the **world origin**. `protr_of` (`run_one.py:108`) takes `norm(pos - pos.mean(0))` — distance from the **tissue centroid**. Both are called “protrusion”. Asymmetric growth and extrusion make the vesicle drift, so the origin-referenced version reads drift as elongation and the tube-clustering directions (`cen[i]/rad[i]`) point at the origin instead of along the tube. The only admitted tube counter, `ta_n_tubes_final`, inherits this.

18.2 Okuda Table 1 versus the search box

Three of his values are *outside* the declared ranges; that alone forecloses a faithful run.

Okuda (Table 1)	value	composition_space.py	status
α , Hill coefficient (growth)	10	<code>alpha</code> (1.0, 8.0), def 4.0	out of range
φ , inhibitor diffusivity	10	<code>d_h</code> (0.1, 2.0)	out of range
χ , spatial characteristics	0.001–0.1	<code>chi</code> (1.0, 10.0)	different quantity
γ , time characteristics	0.01–100	<code>rd_rate</code> (0.2, 3.0)	1.2 of 4 decades
ρ_{sw} , growth switch	0.5	<code>a_sw</code> (0.2, 6.0), def 1.5	ok
κ_s	0.2	<code>kappa_s</code> , def 0.2	matches
η_c	0.25	drag in <code>shape_energy_3d</code>	to verify
τ_{cycle}	$50 \eta_c / \kappa_s = 62.5$	<code>min/max_cycle</code>	needs mapping
initial cells	200 / 2000 / 4000 grown	<code>n_cells</code> (150, 2000)	reservoir > 4000

Two *inequalities* define his regime and should be asserted rather than tuned: $k_v \gg \kappa_s v_{ref}^{2/3}$ (incompressibility) and $\tau_{cycle} \gg \eta_c / \kappa_s$ (quasi-static).

18.3 Phases

Phase 0 — stop the bleeding (half a day). Nothing runs unattended until these land.

1. Install `check_alignment` at all three clamp sites in `turing_vertex_study.py`; replace `min(...)` with an assertion. Decides the phantom question.
2. Fix `Q`: checkpoint the end state and continue from it. Until then remove the `q` term from `score_run` — a constant carrying weight 1.0 is worse than no term at all.
3. Fix `ta_*`: subtract the tissue centroid in `_cell_centroids`, and route `protr_of` through the same helper so the two cannot diverge again.
4. Reorder `escalation.decide` (backlog check before the freeze check) and stop `round.py:344` discarding a Watcher-vetoed parent from the frontier.

Phase 1 — re-measure what Phase 0 invalidated (one wave, ~1 h). Re-run the instrument gate on corrected `diag.json`; re-run the 32-run study at a 10× reservoir with the alignment assertion live. Either OR001 survives as a clean structural result — what the method paper wants — or it dissolves.

Phase 2 — the Okuda-faithful reference (~1 week). Widen `alpha` to include 10, `d_h` to include $\varphi = 10$, `rd_rate` to four decades; expose the d_a/d_h ratio as the named spatial axis in place of the misleading `chi`. Add an `okuda_table1.yaml` fixture asserting κ_s , η_c , τ_{cycle} , ρ_{sw} and the two regime inequalities. Implement OR001 as Okuda’s *published* method (relax-to-residual plus 3D IH/HI reconnection), and add the missing deformation→patterning arrow: flux weighted by shared interface area s_{jk} and volume v_j , where `tyssue_rd_ops.py:207` is today a degree-normalised Laplacian that drops the geometry entirely — the half of the coupling that lets patterning respond to shape.

Phase 3 — the four metrics that make the figure scoreable (2 days). Spot count (Fig. 7), tube diameter (Fig. 5a vs 5b), branch count over time (Fig. 6), curvature-mode spectrum. Calibrate the spatial axis by matching Okuda’s ~5 spots at 2000 cells, then hold it fixed.

Phase 3b — the camera, and why it is a metric problem. `run_one.py:335` hard-codes `azim=30` and `_draw` sets `view_init(elev=18, azim)`, so a tube growing into the far hemisphere hides behind the body. A fixed top view fails the same way — a tube along $+z$ becomes a dot. Instead **lock the camera to the protrusion axis**, which `tube_diameter` already computes: panel 1 perpendicular to it (tube always in full profile), panel 2 looking *down* it (end-on). The end-on view *is* the diameter and branch-count readout, so camera and metric share one axis; a fixed top-down (`elev = 88`) can be a third tile. The same change fixes the montage: with a locked axis, one shared `Lbox` is comparable across runs.

Phase 4 — the demonstration (one wave). The designed (χ, γ) grid: 200 cells for tubulation, 2000 for undulation, six to eight jobs. This is the figure that demonstrates the method.

Phase 5 — the method claim. Re-enter the loop against the now-known-good reference, where ablation and necessity testing are genuinely informative. That, with the pre-registration / surprise

/ retraction record, is the method result — and the order matters: the loop earns its claim by being run against a reference it did not itself produce.