

Plexus

A blueprint for differentiable tissue

Janelia Research Campus

July 19, 2026

Abstract

We introduce *Plexus*, a typed intermediate representation for biological mechanisms. Rather than defining yet another simulator, Plexus defines a common language whose primary objects are biological operators. Every operator specifies the biological entities it acts upon, the states it requires, the relations between these entities, and one or more numerical implementations. The Plexus research program has three objectives. First, to define an algebra of biological operators. Second, to decompose the existing corpus of simulation models into this common language. Third, to construct differentiable implementations of these operators so that mechanistic models can be directly inferred from biological data.

1 Introduction

Computational biology has produced a rich ecosystem of simulation frameworks. Neural simulators model electrical activity, reaction–diffusion systems model chemical transport, particle and continuum methods model mechanics, while vertex, Cellular Potts and active-matter models describe multicellular tissues. These frameworks remain however largely incompatible. Combining electrical signalling, mechanics, metabolism and development often requires coupling multiple specialized simulators whose underlying abstractions differ fundamentally. Each simulator introduces its own representation of biological mechanisms, making models difficult to compare, compose and extend. This fragmentation is not primarily a numerical problem but more a language problem.

Instead of treating simulation frameworks as fundamental, we treat biological mechanisms as the primary objects. To this aim, Plexus introduces an intermediate representation whose central elements are biological operators. Operators describe mechanisms such as adhesion, diffusion, division, chemotaxis, electrical propagation or active stress independently of their numerical implementation. In this perspective, biological modeling becomes analogous to compiler design. Just as modern compilers separate programming languages from their implementations through an intermediate representation, Plexus separates biological mechanisms from their numerical implementations.

This viewpoint defines a broader research program. First, Plexus seeks to construct an algebra of biological operators whose elements can be composed into complex dynamical systems, compared through their observable consequences, and expanded as new mechanisms are discovered. Second, Plexus aims to extract this operator algebra from the existing corpus of biological simulation models. Rather than viewing current frameworks as competing simulators, we regard them as repositories of reusable mechanisms. By decomposing these models into a common language, we seek to identify a compact vocabulary capable of expressing a wide range of biological dynamics across scales. Third, Plexus constructs differentiable counterparts of these operator compositions.

Differentiability allows mechanistic models to be fitted directly to biological data, residuals to be interpreted mechanistically, and missing operators intuited.

Finally, because biological mechanisms are represented in a formal language rather than embedded in simulator-specific code, they become amenable to automated reasoning. The explicit separation between biological semantics and numerical implementations allows agentic systems to operate directly on the language itself, composing, comparing and optimizing mechanistic models without modifying simulator code. The search for biological explanations therefore becomes a search over model structures rather than simulator implementations. In sum, mechanistic discovery can proceed by iteratively assembling operator compositions, fitting them to observations, and extending the language whenever existing operators prove insufficient.

2 The language

Plexus defines a typed language for expressing biological mechanisms independently of their numerical realization. The language is built around four fundamental concepts: *operators*, *sets*, *states*, and *fields*. Biological models are therefore represented as compositions of operators acting on discrete and continuous biological objects rather than simulator-specific algorithms.

2.1 Operators

Operators are the fundamental objects of the language. An operator represents a single biological mechanism, such as adhesion, chemotaxis, diffusion, electrical propagation, active stress, metabolism, growth, or cell division. Every operator is defined by a typed signature and one or more numerical implementations. The signature specifies *what* biological transformation is performed, whereas implementations specify *how* it is numerically realized. The signature specifies

- the input sets,
- the output sets,
- the state variables consumed and produced,
- the maps relating the participating sets.

Maps define the relations required by the mechanism. They determine, for example, which particle belong to which cell, which synapses connect which neurons, or how cells relate to a continuous field. Plexus incorporates maps into the operator, making operators typed morphisms between biological sets.

2.2 Sets

Sets represent collections of biological entities. Examples include genes, proteins, particles, cells, neurons, synapses, tissues, organs and organisms. Sets naturally organize into hierarchies reflecting biological organization. Genes regulate proteins, proteins determine cellular behaviour, cells assemble into tissues, tissues form organs, and organs compose organisms. Each level is represented uniformly as a set, allowing the same language to describe intracellular signalling, multicellular mechanics, developmental processes and organ-level physiology.

2.3 States

Every element of a set carries a state describing its current biological condition. While sets define *what* biological entities exist, states define *what each entity carries*. The state schema is specific to each set. A particle may carry position, velocity and deformation gradient; a neuron membrane voltage and calcium; a cell polarity, stiffness and gene-expression levels; a synapse conductance and plastic weight. Operators transform states. They consume selected state variables and produce updates to others, while leaving the underlying biological entities unchanged.

2.4 Fields

Fields represent continuous quantities defined over space rather than over discrete biological entities. Examples include morphogen concentrations, pressure, extracellular signalling molecules, electric potentials, and material properties. Fields complement sets. While sets describe discrete biological entities, fields describe the continuous environment in which these entities evolve. Operators couple sets and fields through explicit maps relating discrete entities to continuous spatial representations.

3 The operator algebra

The language introduced in the previous section defines the objects manipulated by Plexus. We now define the elementary operator families from which biological models are constructed. The central hypothesis of Plexus is that a large fraction of existing biological simulation models can be expressed as compositions of a small number of elementary operator families. Rather than introducing a new operator for every biological phenomenon, Plexus identifies a compact operator algebra that can be reused across mechanics, signalling, metabolism, development and physiology.

The proposed operator algebra consists of seven elementary families (Fig. 1):

- **Lateral:** interactions within a biological set.
- **Aggregate:** many-to-one transformations across biological scales.
- **Broadcast:** one-to-many transformations across biological scales.
- **Exchange:** transformations between sets and fields.
- **Rewire:** modification of relations between biological entities.
- **Divide:** creation of biological entities.
- **Die:** removal of biological entities.

The first four families transform biological state while preserving the underlying biological organization. Rewire modifies the interaction graph, whereas Divide and Die modify the biological entities themselves. Together they form the elementary computational vocabulary of Plexus.

Lateral. Lateral operators describe interactions between entities belonging to the same set. Typical examples include mechanical forces between particles, adhesion between neighbouring cells, electrical communication between neurons, or biochemical interactions within a metabolic network.

Aggregate. Aggregate operators propagate information from many entities to a higher level of biological organization. Examples include computing cellular properties from their constituent particles or accumulating synaptic currents acting on a neuron.




Figure 1: The elementary operator algebra of Plexus. Lateral, Aggregate, Broadcast and Exchange transform biological state. Rewire modifies biological relations, whereas Divide and Die modify biological entities. Complex biological models are obtained by composing these elementary operator families.

Broadcast. Broadcast operators perform the complementary transformation, propagating information from higher organizational levels toward their constituents. Aggregate and Broadcast therefore define the upward and downward flow of information across biological scales.

Exchange. Exchange operators couple discrete biological entities with continuous fields. Typical examples include secretion, sensing, transport, uptake, or particle-grid transfer.

Rewire. Rewire operators modify the relations between biological entities without changing the entities themselves. They update neighbour graphs, contact networks, developmental connectivity or other interaction structures, determining which entities participate in subsequent operator evaluations.

Divide and Die. Living systems continuously modify their own composition. Divide creates new biological entities, whereas Die removes existing ones. These operators modify the biological organization itself rather than only its state.

4 Composing biological mechanisms

Biological models are expressed as compositions of elementary operators. Rather than embedding all mechanisms within a monolithic simulator, Plexus represents a model as an ordered sequence of operator applications. This sequence, called the *schedule*, specifies both the order and the frequency with which operators are evaluated. Formally, one simulation step is the composition

$$\Phi = \mathcal{O}_n \circ \dots \circ \mathcal{O}_1,$$

whose repeated application generates the system dynamics

$$x_{t+1} = \Phi(x_t).$$

The schedule therefore describes the temporal organization of biological mechanisms. Different operators may evolve on different time scales. Mechanical interactions may be evaluated every simulation step, whereas diffusion, growth or cell division may occur less frequently. A biological model therefore becomes a multi-rate composition of biological mechanisms rather than a monolithic numerical solver.

Operators are pure transformations. Rather than modifying biological state directly, every operator returns its contribution to the system evolution. If operator $\mathcal{O}_i$ produces the increment Δ_i , the complete update is obtained by

$$\Delta = \sum_i \Delta_i,$$

before being passed to the numerical implementation. This formulation has two important consequences. First, biological mechanisms remain independent and composable: introducing a new mechanism simply requires adding a new operator to the schedule. Second, the language remains independent of its numerical realization. The schedule specifies the biological model, whereas the numerical implementation specifies how that model is evaluated.

5 Numerical implementations

The operator algebra defines the biological semantics of a model independently of its numerical realization. Every operator may admit multiple implementations satisfying the same typed signature. These implementations differ in numerical realization while preserving identical biological semantics. Selecting an implementation is therefore a numerical decision rather than a biological one.

This abstraction naturally accommodates a broad range of computational methods. Mechanical operators may be implemented using Material Point Methods, finite elements, vertex models, phase-field methods, Cellular Potts models or optimal-transport formulations. Signalling operators may employ Hodgkin–Huxley equations, graph neural networks or Neural ODEs. Likewise, field operators may use finite differences, finite elements, neural implicit representations or spectral methods. These implementations coexist within the same language because they satisfy identical operator contracts.

Different implementations may also employ different numerical integration schemes. Depending on the operator, implementations may evolve first-order states such as concentrations or membrane voltages, second-order states such as mechanical accelerations, or rely on explicit, implicit or adaptive solvers. These choices affect computational efficiency and numerical accuracy but not the biological interpretation of the model.

This distinction fundamentally changes how biological simulations are compared. Traditional frameworks compare complete simulators, intertwining biological mechanisms with numerical algorithms. Plexus instead separates these two questions. Biological models are compared through the operators they compose, whereas numerical methods are compared through the implementations realizing those operators.

This separation establishes two independent axes of scientific progress. Biological innovation corresponds to discovering new operators or new operator compositions, whereas numerical innovation corresponds to developing more efficient or more accurate implementations of existing operators. Existing simulation frameworks therefore become complementary backends implementing a common intermediate representation rather than competing modelling paradigms.

6 Mechanistic inverse modelling

One of the primary motivations for introducing a formal language of biological mechanisms is to identify mechanistic models directly from experimental observations. Rather than calibrating simulations manually, the objective is to infer the parameters, spatial organization and constitutive laws governing biological systems from imaging data.

Plexus makes this possible by associating every biological operator with one or more differentiable implementations satisfying the same typed signature. Differentiability therefore becomes a property of the language rather than of a particular simulator. Regardless of the numerical backend, gradients propagate through compositions of explicit biological operators rather than through monolithic simulation code.

This distinction fundamentally changes the role of inverse modelling. Traditional approaches assume that the underlying mechanism is already correct and seek the parameter values that best reproduce the observations. Such parameter optimization is effective when the forward model is complete, but it becomes misleading when important biological mechanisms are absent. Optimizing existing parameters may improve the fit while merely compensating for missing mechanisms.

In Plexus, parameters remain attached to explicit biological operators rather than to opaque numerical models. Mechanical properties belong to mechanical operators, reaction rates to biochemical operators, conductances to signalling operators, and spatial fields to the mechanisms generating them. Parameter estimation therefore preserves the mechanistic interpretation of the model.

Residuals likewise acquire a mechanistic interpretation. A persistent mismatch between simulation and experiment no longer indicates only poor parameter values; it may also indicate that the current operator composition is unable to explain the observations. Differentiability thus provides more than efficient optimization: it enables quantitative comparison between competing mechanistic models while preserving their explicit biological structure.

7 Agentic mechanistic discovery

The explicit representation of biological mechanisms transforms scientific discovery from a parameter optimization problem into a language exploration problem. Because mechanisms are represented as explicit operators rather than embedded in simulator-specific code, they become directly accessible to agentic systems.

Plexus proposes a three-loop strategy for mechanistic discovery. Each loop operates at a different level of abstraction, progressively moving from existing models toward new biological knowledge.

Loop I: Understanding. The first loop analyses an existing simulation model. Rather than treating the simulator as opaque code, the agent decomposes it into Plexus operators, identifies the biological mechanisms it implements, and reconstructs an explicit mechanistic description. Existing simulation frameworks therefore become repositories of reusable biological operators.

Loop II: Fitting. The second loop estimates the parameters, spatial fields and initial conditions of the resulting mechanistic model by fitting it directly to experimental observations. Because every operator admits a differentiable implementation, optimization preserves an explicit correspondence between parameters and biological mechanisms.

Loop III: Discovery. The third loop analyses the residual remaining after optimization. Persistent discrepancies are interpreted as evidence that the current operator composition is incomplete rather than simply poorly parameterized. The agent therefore proposes alternative operators, compares competing mechanistic hypotheses, and extends the operator algebra whenever existing mechanisms fail to explain the observations.

Together, these three loops transform simulation, inverse modelling and hypothesis generation into successive stages of a single scientific workflow. The language itself becomes the object of discovery: new biological knowledge is represented by extending the operator algebra rather than merely adjusting parameter values.

A Reference implementation

The main text defines Plexus as a formal language for expressing biological mechanisms. This appendix describes the reference implementation of that language. The implementation follows the abstractions introduced in the paper without introducing additional biological concepts. Its purpose is simply to execute models written in the Plexus language.

The implementation distinguishes two levels. The *language* defines the biological semantics through operators, sets, states and fields. The *runtime* provides concrete software classes implementing these concepts.

Language	Reference implementation
Operator (contract)	<code>OperatorContract</code>
Implementation	an <code>Operator</code> subclass, registered under a contract
Set	<code>Level</code>
State	typed tensors stored by a <code>Level</code> (sliced by its <code>state_schema</code>)
Map	an index buffer on a <code>Level</code> (<code>parent</code> , <code>pre</code> , <code>post</code>)
Field	<code>Field</code>
Biological model	<code>Hierarchy</code>
Registry	<code>@register_operator</code> / <code>@register_entity</code> / <code>@register_field</code>

Maps are not a separate runtime object. Consistently with the operator-centric view of the main text, a map is an index buffer carried by a `Level` (`parent`/`parent_name` for containment, `pre`/`post` for incidence) and *named* by the operators that traverse it in their typed signature — so the runtime realizes four primitives (operators, sets, states, fields), with maps folded into the operator.

A.1 Hierarchy

The `Hierarchy` represents one executable biological model. It is the runtime container responsible for assembling the language primitives into a simulation.

A `Hierarchy` stores

- the collection of biological `Levels`,
- the continuous `Fields`,
- the instantiated `Operators`,

- the execution Schedule,
- backend-specific runtime data.

The Hierarchy therefore corresponds to one complete biological model expressed in the Plexus language.

A.2 Levels

A **Level** is the runtime representation of a biological Set. Its elements are stored as flat batched tensors (never one object per element).

Concretely, a Level carries

- a **state** tensor and the **state_schema** that names its columns (**pos**, **vel**, **voltage**, ...);
- an integer **depth** (a scale hint; 0 for a leaf set);
- optional biological **types** (parametric variants within the set);
- its maps, as index buffers: **parent** (with **parent_name**) for containment, and **pre/post** for the incidence of an edge-set;
- an occupancy vector **occ** $\in [0, 1]$ marking the live subset, so a cardinality-changing operator (Divide/Die) wakes or retires fixed slots without resizing.

State is stored as contiguous tensors for computational efficiency, while its biological interpretation is determined entirely by the **state_schema** declared by the language. Consequently, the runtime never assumes that a particular state variable represents position, voltage, concentration or any other specific quantity: operators read named blocks (`lvl.get("pos")`) rather than hard-coded column offsets.

A.3 Fields

Fields represent continuous spatial quantities. Unlike Levels, which contain discrete biological entities, Fields describe continuous media such as morphogen concentrations, extracellular signals, pressure or material properties.

Operators interact with Fields exclusively through the maps declared by their typed signatures. The numerical realization of these maps (interpolation, particle-grid transfer, neural implicit evaluation, etc.) is entirely implementation dependent.

A.4 Operators

The runtime realizes the contract/implementation split of Section 5 directly. A biological operator is an **OperatorContract** — a fixed record of (**name**, **kind**, **family**, **set**, **signature**) together with a dictionary of interchangeable **implementations** and a **default**. Each implementation is an **Operator** subclass; the contract owns the biology, the subclass owns only the numerics.

An operator is registered by decorating its implementation class. The FIRST registration of a name creates the contract from the class's typed **signature()**; a later registration of the same name with a different **implementation=** adds an interchangeable implementation to that same contract (the engine checks the two share a **kind**, so only the numerics may differ).

Listing 1: An operator implementation: the decorator names the contract (set/kind/family) and this implementation; the class attributes are the typed signature and the capabilities.

```
@register_operator("neuron_voltage", set="neuron", kind="lateral",
                  family="interaction", implementation="hodgkin_huxley")
class HodgkinHuxley(Lateral):
    # typed signature -- the biology (Plexus2 sec. 2.1)
    INPUTS = ["neuron", "synapse"] # input sets
    OUTPUTS = ["neuron"] # output sets
    READS = ["voltage", "w"] # state consumed
    WRITES = ["voltage"] # state produced
    MAPS = ["pre", "post"] # maps traversed (folded into the operator)
    # capabilities -- the numerics
    EMIT = "velocity" # integration order of the returned delta
    SUPPORTED_DIMS = [2, 3]
    DIFFERENTIABLE = True
    def forward(self, H, mask=None):
        ... # returns {"neuron": d_voltage}
```

Three things follow from this shape.

Signature. `signature()` returns the typed morphism `{inputs, outputs, reads, writes, maps}`, defaulting the sets to the acting set `SET` when a class leaves them unset. This is the machine-readable form of “operator as a typed morphism between sets”; the validator and the atlas read it, the engine never does.

Capabilities. Each implementation declares `SUPPORTED_DIMS` and `DIFFERENTIABLE`; the contract’s `capabilities()` tabulates them across implementations, so an inverse loop can select, e.g., the differentiable implementation of a contract.

Kind and purity. Every operator belongs to exactly one `KIND` — the seven of the main text (`lateral`, `aggregate`, `broadcast`, `exchange`, `field`, `rewire`, `structural`). Dynamics operators are pure: `forward` returns a per-set delta and the engine integrates once per tick (its order set by `EMIT` $\in$ `{velocity, acceleration, mpm.acceleration}`); `field/rewire/structural` operators mutate the field, the relation, or the membership in place and return nothing.

Selection is by name: `get_operator(name, implementation)` returns the chosen class (the default when none is named), and `get_contract(name)` exposes the whole contract. The registry therefore separates biological semantics from numerical realization: new numerical methods are new registered implementations, never edits to the execution engine.

A.5 Schedules

A biological model is executed by evaluating an ordered Schedule of operator applications.

During each simulation step

1. operators are evaluated in Schedule order,
2. each dynamics operator returns a per-set delta (or a field/relation/structural operator mutates the grid, the edge set, or the membership in place),
3. the engine sums the deltas acting on each set,
4. each set is integrated once, in the order its operators declare (`EMIT`: first-order `velocity` or second-order `acceleration`).

Because operators are pure and nothing is written in place except the explicitly allowed buffers, gradients flow through the sums and the integrator, so the whole rollout stays differentiable. The Schedule therefore implements the operational semantics of the main text while remaining independent of the particular numerical implementations it composes.

A.6 Model specification

Biological models are described declaratively rather than procedurally, as a single `spec.yaml` with blocks `general` / `sets` / `fields` / `operators` / `schedule`. An operator line names the contract (`op`) and, optionally, which implementation to run (`implementation:`); omitting it selects the contract’s default, so a single-implementation spec is written exactly as before.

A specification therefore defines

- biological Levels and their state schemas,
- Fields,
- the operators to instantiate, each with its selected implementation,
- the execution Schedule.

Before execution the specification is validated against the registered contracts: an unknown operator, an unknown `implementation` for a known operator, a missing required parameter, a selector on an undeclared set/field, or an unresolved schedule token is rejected up front — so a spec that loads is runnable, and the capability checks (e.g. `SUPPORTED_DIMS`) run against the implementation that will actually execute.

A.7 Reference architecture

The reference implementation follows a strict separation of concerns.

- The language (`sets` / `states` / `fields` / `operators`) defines biological semantics.
- The registry stores each operator as an `OperatorContract` with its implementations.
- Implementations (`Operator` subclasses) realize those contracts numerically.
- The validator (`schema.py`) admits only well-formed, runnable specs.
- The execution engine (`engine.py`) builds the `Hierarchy` and iterates the Schedule; it never contains mechanism-specific code.

A single audit (`tools/audit_operator_registry.py`) enforces the contract across the whole registry — every operator has a valid `family` and `kind`, a canonical `EMIT`, and a `SUPPORTED_DIMS` $\subseteq \{2,3\}$ — so the vocabulary cannot silently drift. Consequently, extending Plexus never requires modifying the execution engine: new biological mechanisms are new operator contracts, and new numerical methods are additional implementations of existing operators.

The reference implementation therefore mirrors the architecture proposed in the main text: the language remains stable while implementations and numerical backends evolve independently.

B Building the Plexus operator atlas

One of the principal objectives of the Plexus research program is to construct a comprehensive atlas of biological operators by systematically analysing the existing corpus of biological simulation models. Rather than cataloguing simulators, the atlas catalogues reusable biological mechanisms together with their numerical implementations.

This atlas serves two complementary purposes. First, it provides the biological vocabulary from which new models are assembled. Second, it defines the search space explored by the agentic discovery framework described in the main text.

B.1 From simulators to operators

Existing simulation frameworks are not imported directly into Plexus. Instead, each framework is systematically decomposed into the language introduced in this paper.

The decomposition proceeds through the following steps:

1. identify the biological sets,
2. identify the associated state variables,
3. identify continuous fields,
4. identify biological operators,
5. identify the maps required by each operator,
6. separate biological semantics from numerical implementations,
7. register each operator in the Plexus language.

The result is an implementation-independent description expressed entirely in terms of operators, sets, states and fields.

B.2 Normalization

Different simulators frequently implement identical biological mechanisms using very different numerical algorithms. The atlas therefore groups operators according to biological semantics rather than implementation.

For example,

Biological operator	Possible implementations	Example frameworks
Mechanical interaction	MPM, FEM, Vertex, Phase Field	Taichi, Chaste, Tissue Forge
Chemotaxis	Finite differences, Neural fields	Morpheus, CompuCell3D
Electrical propagation	Hodgkin–Huxley, Neural ODE, GNN	NEURON, Jaxley, BrainPy
Diffusion	Finite differences, FEM	Morpheus, VirtualLeaf
Cell division	Vertex, CPM, Agent-based	Chaste, CompuCell3D

The objective is therefore not to catalogue implementations but to identify the smallest reusable vocabulary capable of reconstructing the existing simulation literature.

B.3 Operator metadata

Every entry in the atlas records

- the operator name, family and kind,
- its typed signature (input/output sets, state read/written, maps),
- its implementations and their capabilities (supported dimensions, differentiability),
- its biological interpretation (`MECHANISM_TAGS`) and the role of each tunable parameter (`PARAM_ROLES`),
- originating publications and reference repositories,
- validation status.

The current registry already carries most of this per operator — name, family, kind, acting set, typed signature, implementations, capabilities, mechanism tags and parameter roles, with provenance kept in the operator’s docstring; the atlas adds the cross-framework publication and repository record. Consequently the atlas simultaneously documents biological mechanisms, available numerical realizations, and scientific provenance.

B.4 Atlas growth

The atlas is intended to evolve continuously.

Every newly analysed simulator contributes either

1. an implementation of an existing operator,
2. a refinement of an existing typed signature,
3. or a genuinely new biological operator.

The first outcome strengthens numerical diversity. The second improves the language. The third expands the biological vocabulary itself.

B.5 Validation of the operator algebra

The atlas provides a direct experimental test of the central hypothesis of Plexus.

If the decomposition of a sufficiently broad collection of biological simulation frameworks converges toward a compact and reusable operator vocabulary, then the proposed operator algebra constitutes a meaningful intermediate representation for biological modelling.

Conversely, repeated discovery of genuinely new operator families indicates that the language remains incomplete and must be extended.

The operator atlas therefore serves both as a repository of biological knowledge and as the principal mechanism through which the Plexus language itself evolves.

C Implementing new operators

The Plexus language grows by introducing new biological operators rather than by modifying the execution engine. A new operator should be introduced only when an existing operator cannot

express the mechanism without changing its biological meaning. New solvers, discretizations, parameterizations or hardware targets should instead be added as alternative implementations of an existing operator.

C.1 Define the biological mechanism

Before implementation, the mechanism must be described independently of its numerical realization. This description should specify

- the biological transformation represented by the operator,
- its operator family,
- its input and output sets or fields,
- the state variables it reads and produces,
- the maps required to relate the participating objects.

Together, these elements define the operator’s typed signature, which serves as its biological contract.

C.2 Separate semantics from implementation

The operator contract and its numerical implementations are represented separately.

A single biological operator may admit multiple implementations based, for example, on Material Point Methods, finite elements, vertex models, phase fields, Neural ODEs or graph neural networks. These implementations may differ in their numerical assumptions, spatial representation, supported dimensions or differentiability while preserving the same biological semantics.

Conversely, mechanisms should not be merged merely because they share similar code or numerical utilities. Distinct biological transformations require distinct operator contracts.

C.3 Register the operator and its implementations

Once defined, the operator is added to the registry by decorating its implementation class with `@register_operator(name, set=, kind=, family=, implementation=)` (Listing 1). The decorator creates the contract on the first implementation of a name and records

- the canonical operator name (plus any aliases), family and kind,
- its typed signature, declared as the class attributes `INPUTS / OUTPUTS / READS / WRITES / MAPS`,
- the available implementations, keyed by `implementation` name,
- their capabilities (`SUPPORTED_DIMS`, `DIFFERENTIABLE`),
- their scientific provenance (in the class docstring).

A second numerical method for the same mechanism reuses the same `name` with a new `implementation=` and is checked to share the contract’s `kind`. Registration makes the mechanism available to every compatible Plexus model without introducing biological knowledge into the execution engine.

C.4 Validation

Every implementation should satisfy four requirements:

1. it satisfies the declared typed signature;
2. it reproduces the intended biological mechanism;
3. its numerical behaviour is verified;
4. its scientific provenance is preserved.

In the reference implementation the first three are gated concretely: the registry audit (`audit_operator_registry`) checks the declared contract, schema validation checks that a spec using the operator loads, and a minimal live spec plus the test suite check its numerical behaviour.

Validation establishes structural and numerical consistency, but does not by itself establish biological validity. The latter must be assessed through comparison with observations, known limits, ablations or the originating model.

C.5 Promotion

Prototype implementations should not enter the core registry directly. Promotion requires a stable operator contract, at least one validated implementation, clear scientific provenance, and evidence that the mechanism is reusable beyond its originating prototype.

The guiding principle is

Prototype freely; promote biological mechanisms conservatively.

D Refactoring guidelines

The current Plexus repository predates the formalization of the language presented in this paper. Consequently, much of the biological semantics is embedded directly in numerical implementations. Refactoring aims to recover these semantics and progressively reorganize the codebase around the language abstractions introduced in this work.

D.1 General principles

Every refactoring should follow the same principles.

- Preserve biological behaviour.
- Improve separation between biological semantics and numerical methods.
- Replace implicit assumptions by explicit typed signatures.
- Prefer extending the language rather than adding special cases.
- Refactor incrementally while preserving existing functionality.

D.2 Recommended workflow

When refactoring an existing module, contributors should proceed in the following order.

1. Identify the biological mechanism implemented by the code.

2. Determine whether it corresponds to an existing operator.
3. If not, introduce a new operator contract.
4. Separate the biological contract from the numerical implementation.
5. Register the implementation.
6. Validate that the refactored implementation reproduces the original behaviour.

D.3 Code migration

Existing code should progressively migrate toward the following organization.

- Biological concepts belong to operator contracts.
- Numerical algorithms belong to implementations.
- Execution logic belongs to the runtime.
- Data structures belong to Levels, Fields and Hierarchies.

The execution engine should never contain mechanism-specific code.

D.4 Incremental evolution

Refactoring is expected to be continuous. During the transition, legacy modules and language-native modules may coexist. New developments should preferentially follow the language architecture, while legacy components are progressively decomposed into reusable operators as they are revisited.

The objective is not to rewrite the repository, but to allow the language to emerge progressively from the existing implementation.

E Reference corpus for the operator atlas

The Plexus operator atlas is constructed by systematically decomposing existing simulation frameworks into biological operators, typed signatures and numerical implementations. To ensure reproducibility, the initial atlas is built from open-source projects for which both the scientific publication and an implementation are publicly available.

Rather than treating each simulator as an indivisible software package, every framework is analysed according to the methodology described in Appendix B. The objective is to recover the reusable biological mechanisms implemented by each project and register them as operators within the Plexus language.

The reference corpus is organized into six broad categories.

E.1 Cell and tissue mechanics

These frameworks describe multicellular mechanics, growth, adhesion, proliferation, remodelling and morphogenesis.

- Chaste

- CompuCell3D
- Morpheus
- Tissue Forge
- VirtualLeaf
- PhysiCell
- Active Vertex Model for Cell-Resolution Description of Epithelial Tissue Mechanics
- Multicellular Simulations with Shape and Volume Constraints using Optimal Transport
- Cell Simulation as Cell Segmentation

Typical operators extracted from these models include

- mechanical interaction,
- adhesion,
- active stress,
- cell growth,
- cell division,
- apoptosis,
- neighbour rewiring,
- particle–cell aggregation,
- optimal transport.

E.2 Reaction–diffusion and signalling

These models describe transport, biochemical signalling and spatial pattern formation.

Representative projects include

- Morpheus
- Reaction-Diffusion Systems in Intracellular Transport and Control
- A Programmable Reaction-Diffusion System for Spatiotemporal Cell Signalling Circuit Design
- SPDEs

Representative operators include

- diffusion,
- reaction,
- secretion,
- uptake,
- field sampling,

- broadcast signalling.

E.3 Neural circuits

These frameworks model neuronal signalling and connectome dynamics.

Representative projects include

- Jaxley
- BrainPy
- Connectome-GNN

Typical operators include

- synaptic transmission,
- membrane integration,
- plasticity,
- neuromodulation.

E.4 Mechanobiology and collective behaviour

Several recent frameworks couple mechanics with biochemical signalling and collective migration.

Representative examples include

- ERK-mediated mechanochemical wave models
- Active Vertex mechanics
- Collective cell migration models
- Bioelectric morphogenesis

Representative operators include

- mechanochemical coupling,
- force transmission,
- wave propagation,
- chemotaxis,
- collective polarity.

E.5 Inverse modelling and differentiable simulation

These projects are particularly relevant for the mechanistic inverse modelling described in Section 6 because they expose differentiable representations of biological systems.

Representative projects include

- MultiCell: Geometric Learning in Multicellular Development
- Differentiable Rendering for 3D Fluorescence Microscopy

- Inferring Cell Junction Tension and Pressure from Cell Geometry
- Machine Learning Interpretable Models of Cell Mechanics from Protein Images
- 3D Single-Cell Shape Analysis using Geometric Deep Learning

Typical operators include

- image formation,
- geometric inference,
- parameter estimation,
- mechanical inversion,
- differentiable observation.

E.6 Image analysis and quantitative biology

These projects provide computational operators that transform experimental observations into structured biological measurements.

Representative projects include

- TissueMiner
- ERNet
- DeXtrusion
- FlowShape
- CellRank

Representative operators include

- segmentation,
- tracking,
- graph construction,
- feature extraction,
- trajectory inference.

E.7 Building the atlas

Each framework in the reference corpus is analysed using the same procedure.

1. Identify the biological sets.
2. Identify their state variables.
3. Identify continuous fields.
4. Identify the biological operators.
5. Identify the maps relating the participating objects.

6. Separate biological semantics from numerical implementation.
7. Register the resulting operators in the Plexus atlas.

The objective is not to reproduce the original software architecture, but to recover the reusable biological mechanisms that underlie existing simulators. Each analysed framework may contribute additional implementations of existing operators, refine existing typed signatures, or reveal genuinely new biological operators. Consequently, the atlas evolves continuously as new open-source simulation frameworks become available.