

Coupling discrete Turing reaction–diffusion to a 3D vertex model: what transfers cleanly, and what needs force balance

Notes from a `plexus2` prototype of Okuda et al. (2018)

July 2026

Abstract

We report the lessons from building a `plexus2` prototype that couples a *discrete* Turing reaction–diffusion (RD) system to a 3D vertex / self-propelled-Voronoi model on a growing epithelial vesicle, following Okuda *et al.*, *Sci. Rep.* **8**:2386 (2018). Four findings stand out. **(1)** When growth is morphogen-*independent*, the RD is a passive *colouring* of the cells: it never touches the mechanics. Apparent “noise” in Gray–Scott patterns was therefore not a chemistry problem but a *timestep* artefact — the coarse step Gray–Scott needs to develop ($\Delta t = 0.35$) destabilises the vertex mechanics, which had to be detuned to stay stable. Running the mechanics at a fine, stable step ($\Delta t = 0.02$) while *time-rescaling* the RD by a single rate multiplier recovers a clean epithelium carrying the coral / labyrinth / holes patterns (cell-shape roughness 0.04 vs. 0.24). This is exactly the chemistry-vs-mechanics timescale separation ($\Delta t_m \ll \Delta t_v$) of the original model. **(2)** A latent operator bug — a missing `SphericalVoronoi` import silently swallowed by a broad `except` — had made the tangential (Lloyd) relaxation a *total no-op* for the operator’s entire life; the cell uniformity seen in earlier renders came *entirely* from a render-side cosmetic, which masked the dead physics. **(3)** Reproducing morphogen-driven *tubulation* (their Fig. 5) is qualitatively harder: an overdamped, single-step vertex relaxation *without* force-balance iteration or topological reconnection converts localised growth into *global* chaotic buckling, uncorrelated with the morphogen ($\text{corr}(a, r) \approx 0$), rather than local tubes. **(4)** A working alternative sidesteps the fragile shell mechanics entirely: amplifying the (already present, but weak) activator-dependent *division* into a sharp $\sim 4:1$ ON/OFF contrast on the fast, stable spherical mechanics produces *localised* extrusion that tracks the morphogen ($\text{corr}(a, r) = 0.31$ vs. ≈ 0 for the shell), while the monolayer stays clean – the high-activator domains bulge out into lobes.

1 The model in one paragraph

One `cell` set carries three independently integrated state blocks: position `pos` (the vertex coordinate), morphogens `chem` = $[a, h]$ (a Gray–Scott or Brusselator RD on the cell–cell adjacency graph, so mechanical deformation reshapes the diffusion), and a target size `v0` that a Hill-saturating growth law inflates. Division wakes a dormant buffer slot when `v0` reaches its threshold. Mechanics is an overdamped shape energy $U = \sum_j K_V (V_j - V_j^0)^2 + K_S S_j$ over ghost-bounded Voronoi cells; the force is $-\nabla U$ by autodiff through the tetrahedral circumcentres. Each of these is a `plexus2` operator (a typed morphism on the set) composed by a schedule; different operators run at different cadences, which is how the multiple timescales are expressed.

2 What transferred cleanly: Fig. 4 (patterns on a growing vesicle)

In Fig. 4 growth is *activator-independent* (homogeneous), so the RD only colours the cells. The naive Gray–Scott vesicle looked badly corrupted — a lumpy, non-epithelial surface — and it was tempting to blame the reaction–diffusion. It was not: the mechanics was simply being integrated at Gray–Scott’s coarse $\Delta t = 0.35$, far outside the vertex model’s stability window. The fix is to *decouple the timescales*: keep the mechanics at the fine step that makes it clean ($\Delta t = 0.02$, the same parameters that produce a nice Brusselator vesicle) and speed the RD up by a rate multiplier ρ (with the diffusion coefficient scaled to match), so the identical pattern develops in the same number of frames. Roughly $\rho \approx \Delta t_{\text{coarse}} / \Delta t_{\text{fine}}$; the only ceiling is the explicit-diffusion CFL limit ($\Delta t d_h \chi \bar{k} \lesssim 0.4$), which the delicate low-feed labyrinth regime hits first and which must be respected per regime. The result (Fig. 1) is a uniform epithelium carrying the branching coral pattern, with shell roughness 0.04 against 0.24 for the coarse-step version. The take-away is general: *when a signal does not feed back on shape, do not let its integration timescale dictate the mechanics*.




Figure 1: **Gray–Scott coral on the vesicle mechanics** (fig4.coral_v), at 0/50/100% of the run. Top: closed external view; bottom: cross-section through the monolayer and lumen. Cells are uniform (near-hexagonal) and the shell is smooth; the reaction–diffusion is a pure colouring of a clean epithelium. $N:600 \rightarrow 3000$, roughness 0.04.

3 A cautionary tale: the silent no-op

The uniform cells above depend on a tangential (Lloyd) relaxation that moves each cell toward its spherical-Voronoi centroid. That operator had *never once* executed: the module imported Delaunay, ConvexHull, cKDTree but not SphericalVoronoi, so every call raised NameError,

which fell into a blanket `except Exception: return zeros`. The knob was inert — `k_lloyd` 0.3/1.0/2.0 gave byte-identical output. The uniformity previously visible came entirely from a *render-side* Lloyd pass (a cosmetic applied only when drawing). The moment we honoured “uniformity should come from physics, not the render” and drew the true cells, the dead operator was exposed. Two lessons: a broad `except` around a geometry call can hide a programming error for the lifetime of a component; and a cosmetic post-process can mask a missing mechanism indefinitely. The fix was one import plus narrowing the handler to the genuine Qhull degeneracies; the operator was also made timestep-independent (move a fixed *fraction* toward the centroid per tick, capped at $0.7\times$ the cell spacing).

4 What did not transfer: Fig. 5 (tubulation)

Fig. 5 makes the activator a *mitogen*: high-activator cells grow, and on a closed sheet the locally excess area must buckle out of plane into tubes — no explicit extrusion force. We implemented the coupling (a Hill switch from activator concentration to target-volume growth) on the apical/basal “shell” mechanics, and swept ~ 15 configurations of domain size, growth rate, bending, surface tension, and lumen pressure. *None* produced clean tubes. The behaviour is bimodal: gentle growth or stiff bending gives a perfectly smooth sphere (no buckling), while enough growth gives *global* chaotic buckling that tears the monolayer into a blob. Decisively, the correlation between activator concentration and radial protrusion is ≈ 0 in every case: the bulges are *not* where the activator is — it is generic Euler buckling of an over-grown, lumen-pressurised shell, not localised extrusion.

Three mechanics ingredients are missing, all present in the original: (i) *force-balance iteration* — Okuda solves the overdamped balance to residual $< 10^{-5}$ each Δt_v , whereas a single explicit step lets individual cells spike rather than forming coherent folds; (ii) *topological reconnection* (RNR / T1), without which cells jam instead of flowing into a tube; and (iii) a tractable per-cell tessellation — our `ConvexHull`-per-cell shell becomes pathologically slow, even hanging, exactly in the heavily-deformed regime tubes require. Clean tubulation is thus gated on a vectorised, force-balance-iterated shell mechanics plus a reconnection operator, not on more parameter tuning.

5 Field notes: what we learned about the model

Grouped by subsystem. The numbers are the values that actually worked — or failed — in this prototype; they are the hard-won part and the main reason for this document.

Timescales — the master knob.

- Gray–Scott needs a large step ($\Delta t \approx 0.35$) to develop in $\sim 10^4$ frames; the overdamped vertex model is stable only for $\Delta t \lesssim 0.05$. Irreconcilable at a single step.
- Resolution: run mechanics at $\Delta t = 0.02$ and multiply the RD rate by $\rho \approx \Delta t_{\text{coarse}} / \Delta t_{\text{fine}} \approx 17.5$, scaling the diffusion prefactor χ by the *same* ρ so the reaction:diffusion ratio — and thus the pattern — is preserved. A pure time-rescaling: identical pattern, stable mechanics.
- Explicit graph-diffusion is stable only while $\Delta t d_h \chi \langle k \rangle \lesssim 0.4$. With $d_h = 0.16$, $\langle k \rangle \approx 14$, $\Delta t = 0.02$ this caps $\chi \lesssim 9$. Robust high-feed coral ($F = 0.058$) tolerates $\chi = 9.8, \rho = 35$; the delicate low-feed labyrinth ($F = 0.029$) *diverges to NaN* there and needs $\chi = 5.6, \rho = 20$ over more frames. Tune the rescaling *per regime*.

Tangential (Lloyd) relaxation.

- Make it timestep-independent: move a fixed fraction `k_lloyd` of the way to the spherical-Voronoi centroid per tick, then divide by Δt to undo the engine's $x += \Delta t v$. `k_lloyd` is then a clean fraction (1.0 = one full Lloyd step/tick; 1.5 slightly overshoots and converges faster). Cap the per-tick move at $0.7 \times$ mean spacing so a defect cell cannot jump neighbours and tangle the mesh.
- It uses `SphericalVoronoi` of the cell *direction* vectors, which degenerates once the surface is non-spherical — so it is valid on the vesicle, *not* on Fig-5 tubes.

Morphogen $\rightarrow$ growth (the coupling).

- $\lambda = \lambda_{\text{ref}}(\rho + g \text{Hill}(a))$, $\text{Hill}(a) = a^\alpha / (a^\alpha + a_{\text{sw}}^\alpha)$. ρ is the OFF multiple, $\rho + g$ the ON one.
- a_{sw} *must* sit inside the activator's actual range or the switch never fires. The Gray–Scott activator peaks near 0.4, so $a_{\text{sw}} = 1.0$ gives $\text{Hill} \approx 0.16$ (a feeble 16% contrast); $a_{\text{sw}} \approx 0.2$ makes it bite. $\alpha = 4$ gives a near-step switch.
- Even “homogeneous” Fig-4 growth ($\rho = 1$) carries the Hill term, so high-activator domains already bulge weakly. Amplify with g ($\rho = 0.5, g = 1.5 \Rightarrow 4:1$) for real localised proliferation.

Division.

- `divide_2x` fires when v^0 reaches `ratioxbase` ($v_{\text{th}} = \frac{4}{3}v_{\text{ref}}$); the daughter wakes a dormant buffer slot, both reset to base, and the daughter inherits the mother's `chem`. Faster growth $\Rightarrow$ earlier division $\Rightarrow$ local density $\Rightarrow$ local area excess — the buckling driver.
- `reset_noise` spreads daughter target sizes so the cycle stays asynchronous (no division waves). Fixed-buffer occupancy: division stops when the buffer is full.

Mechanics — two flavours; know which.

- *Spherical-ghost* `voronoi_tension_3d`: fast (vectorised, $\sim 32 \times$ over the per-cell loop), stable, keeps a clean monolayer — but the global ghost sphere *fights* local deformation, so at most mild bumps. Right for Fig 4 and the differential-growth route.
- *Per-cell apical/basal shell* `voronoi_tension_shell`: allows local buckling but needs a `scipy ConvexHull` per cell; under heavy deformation this is pathologically slow and can *hang* (observed 77-min hang, 84-min no-finish). Not viable for tubes un-vectorised.
- Stiff volume vs. surface, $K_V/K_S \approx 10$ (incompressible cells). Bending k_{bend} HIGH (0.6–0.8) forbids out-of-plane curvature $\rightarrow$ no buckle; $k_{\text{bend}} \approx 0.45$ permits it.

Fig-5 buckling phenomenology (all negative, all informative).

- Bimodal: gentle growth / stiff bending $\rightarrow$ smooth sphere; enough growth $\rightarrow$ global chaotic buckle that tears the monolayer. No clean-tube window in between.
- `lumen_pressure` holds the closed vesicle rigid, so area is stored until it releases as one global Euler buckle at a random site — morphology decoupled from the pattern. Removing the lumen did *not* localise it ($\text{corr}(a, r)$ stayed ≈ 0).

Diagnostics that earned their keep.

- *roughness* = std/mean of cell radii (0.04 clean, 0.24 corrupted); % *hexagons* (closed sphere tops out $\sim 70\%$; a bounded 2D disc lower, its boundary ring never hexagonal); *PROTR* = 99th-percentile/median radius (≈ 1 sphere, > 1.4 protrusions); $\text{corr}(a, r)$ = the decisive extrusion-is-localised test.

Performance / engineering.

- Multi-rate composition via an engine `every:K` gate plus an intra-operator topology cache (`retess_every`): mechanics reuses the tessellation for K ticks, pattern preserved, $\sim 1.7\times$ at $K=10$.
- `torch.compile` gives $\sim 50\times$ on the fixed- N energy but recompiles when N grows (division) — keep it off for growing runs. CPU/scipy is the real bottleneck; on the cluster, one dedicated node per generation ($\sim 3\times$ over contended local runs).

6 Discussion: a cleaner route to extrusion

The negative result above is informative. Localised growth on a *closed, pressurised* vesicle tends to store area until it releases as a single global buckle at a random site — morphology decoupled from the pattern. A better substrate for *local* extrusion is the fast, stable *spherical* mechanics used for Fig. 4 (it never breaks the monolayer), driven by *localised proliferation* rather than global inflation.

Strikingly, that mechanism is already latent in the Fig. 4 runs. Even “homogeneous” growth carries a Hill term, $\dot{v}^0 = \lambda_{\text{ref}}(\rho + \text{Hill}(a))$, so high-activator cells already divide somewhat faster and their domains bulge slightly outward — the small extrusions one can see in the coral cross-section. They are weak only because the switch was mis-set ($a_{\text{sw}}=1.0$ while the Gray–Scott activator peaks near 0.4, so $\text{Hill} \approx 0.16$). Introducing a gain, $\dot{v}^0 = \lambda_{\text{ref}}(\rho + g \text{Hill}(a))$ with $\rho=0.5$, $g=1.5$ and a_{sw} moved into the activator’s actual range, turns this into a sharp $2\times$ -ON / $0.5\times$ -OFF (4:1) *division* contrast that concentrates proliferation — and hence out-of-plane buckling — at the morphogen domains, on mechanics that stay stable and cheap.

This works (Fig. 2). The high-activator domains bulge outward into coherent lobes, the monolayer stays intact (roughness 0.058, cells 54% hexagonal), and crucially the deformation now *tracks the signal*: $\text{corr}(\text{activator}, \text{radius}) = 0.31$, against ≈ 0 for every shell-mechanics attempt. It is not yet the fully elongated tubes of the original — for that the force-balance and reconnection ingredients of §5 are still needed — but it is genuine, localised, morphogen-driven extrusion on a substrate that never breaks. This reframes tubulation as *patterned proliferation* on a robust substrate rather than differential growth on a fragile deformable shell.

Two methodological points recur. First, *signalling is geometry-coupled but not mechanics-coupled unless you say so*: keeping that boundary explicit turned an intractable “noisy patterning” problem into a one-line timescale fix. Second, in a typed-operator language the right unit of change is often a new *implementation* or *parameter* of an existing contract (the 4:1 growth is still the **growth** operator), not a new operator — “prototype freely, promote mechanisms conservatively.”




Figure 2: **Locally-favoured extrusion** (`fig4.coral_ext`): the same coral RD on the fast spherical mechanics, but with a 4:1 activator-driven division contrast ($\rho=0.5, g=1.5, a_{\text{sw}}=0.2$). The high-activator (dark) domains extrude outward into lobes while low-activator regions stay flat; the cross-section (bottom) shows the monolayer lobing out where the signal is high, and remaining a clean single sheet. $\text{corr}(a, r)=0.31$.